

Mining Social Network Graphs

Analysis of Large Graphs:
Community Detection

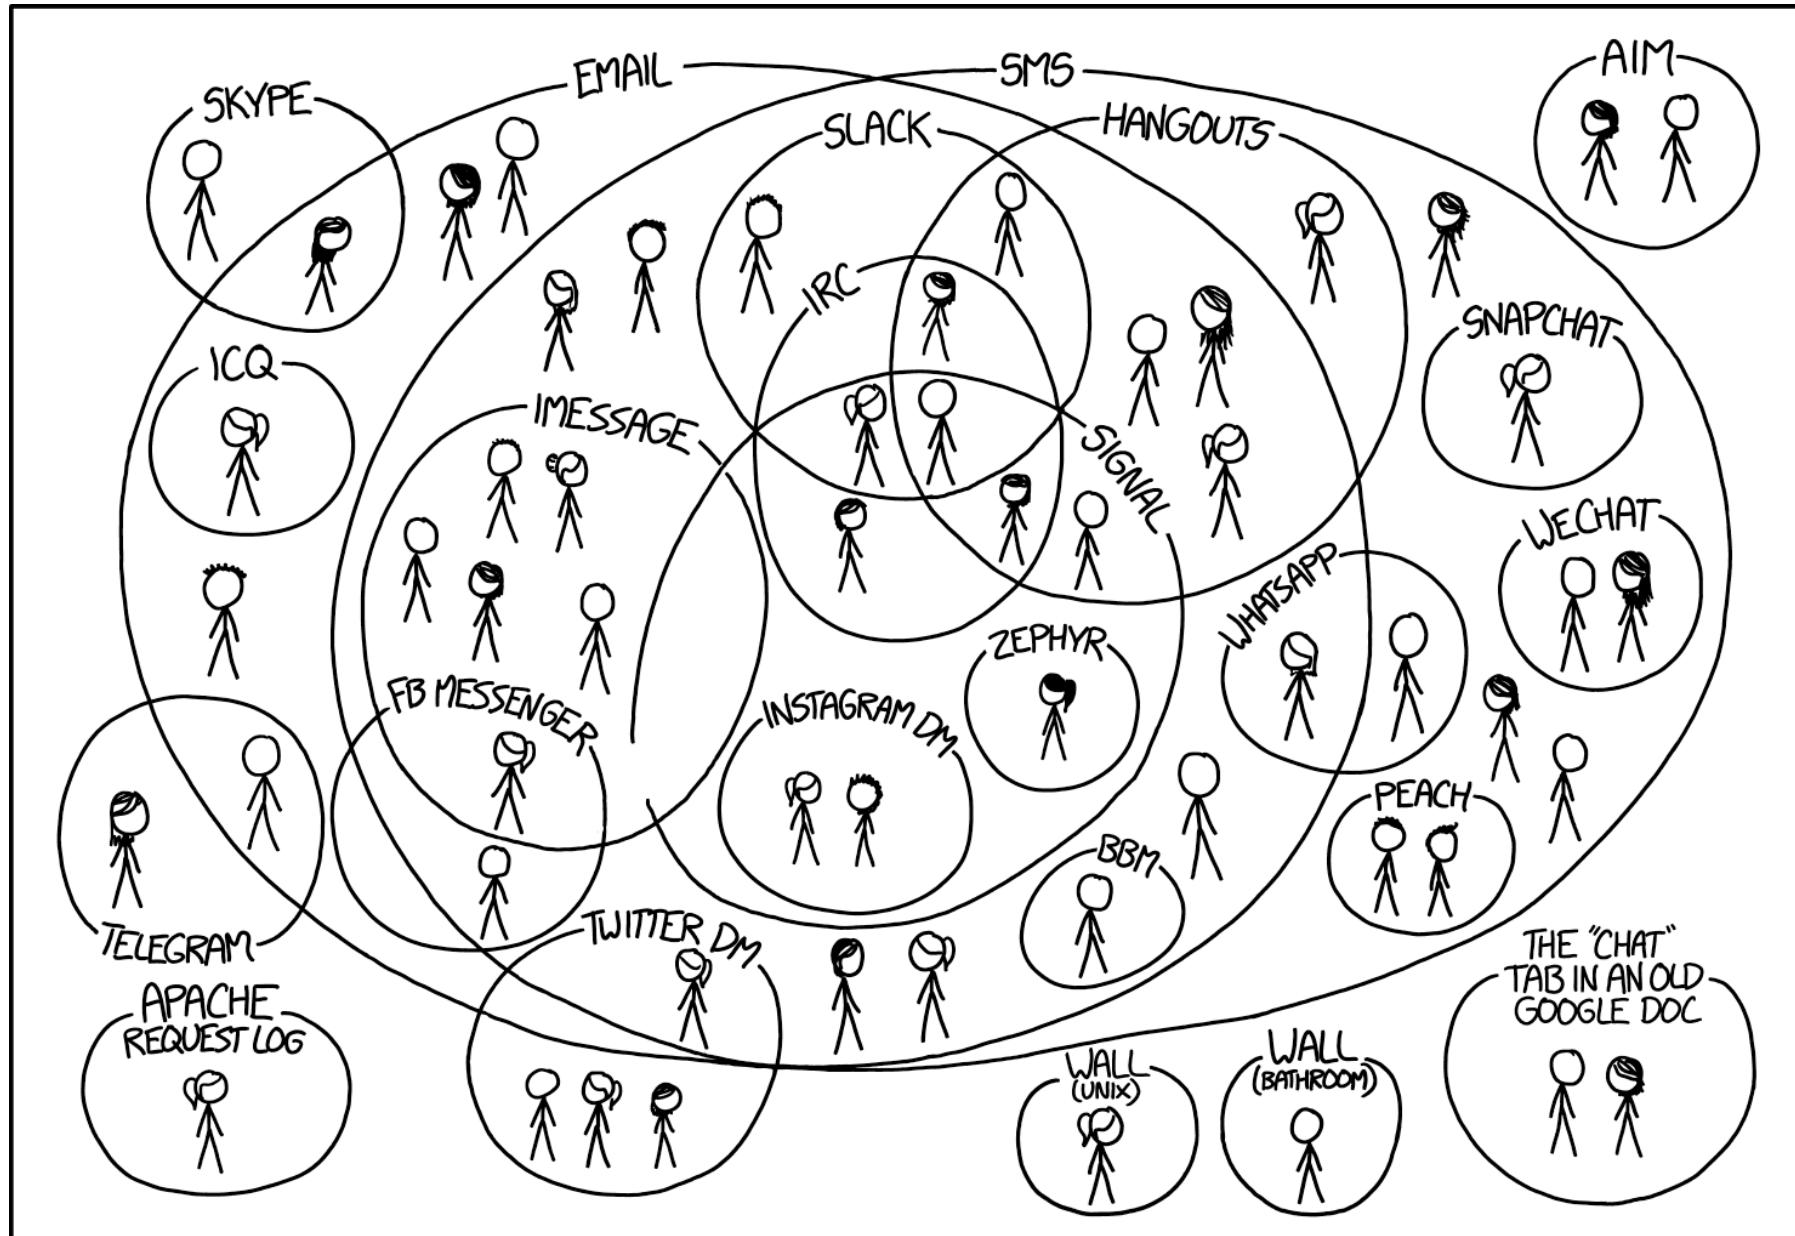
Rafael Ferreira da Silva

rafsilva@isi.edu

<http://rafaelsilva.com>

Note to other teachers and users of these slides: We would be delighted if you found this material useful in giving your own lectures. Feel free to use these slides verbatim, or to modify them to fit your own needs. If you make use of a significant portion of these slides in your own lecture, please include this message, or a link to our web site: <http://www.mmds.org>

With slide contributions from
J. Leskovec, Anand Rajaraman, Jeffrey D. Ullman



I HAVE A HARD TIME KEEPING TRACK OF WHICH CONTACTS USE WHICH CHAT SYSTEMS.

<https://xkcd.com/1810/>

Social Networks: Examples

- Examples

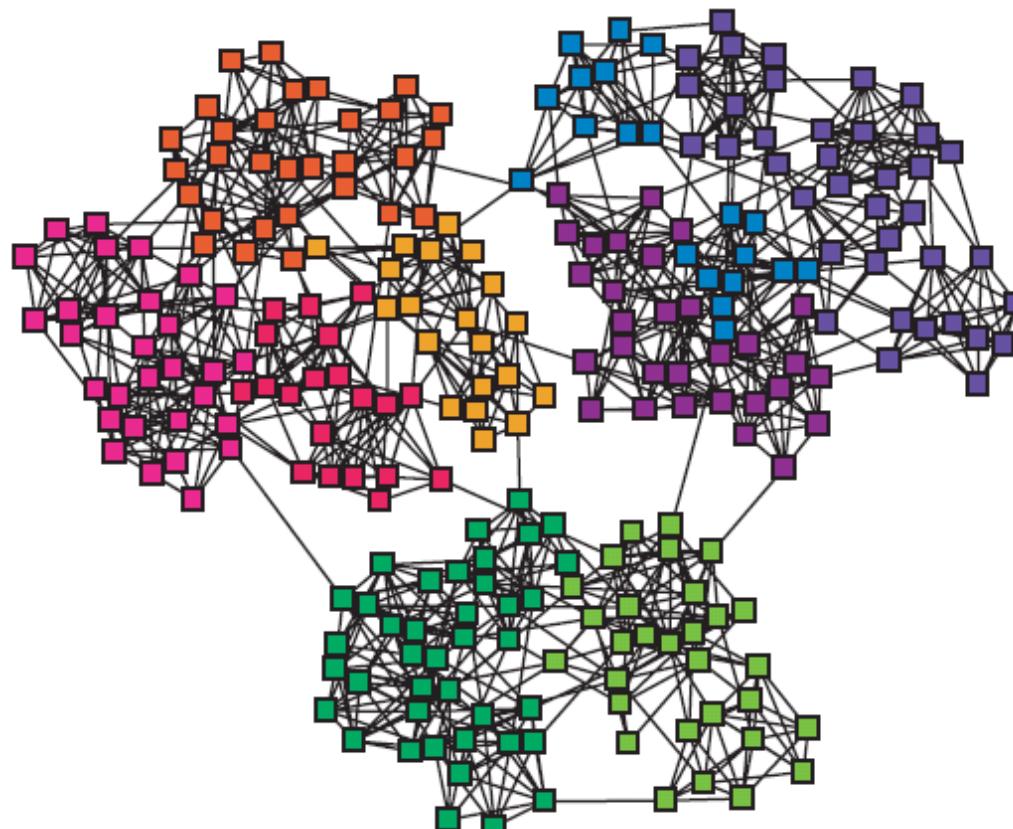
```
docker run -p 8888:8888 jupyter/scipy-notebook:2c80cf3537ca
```

- Use files:

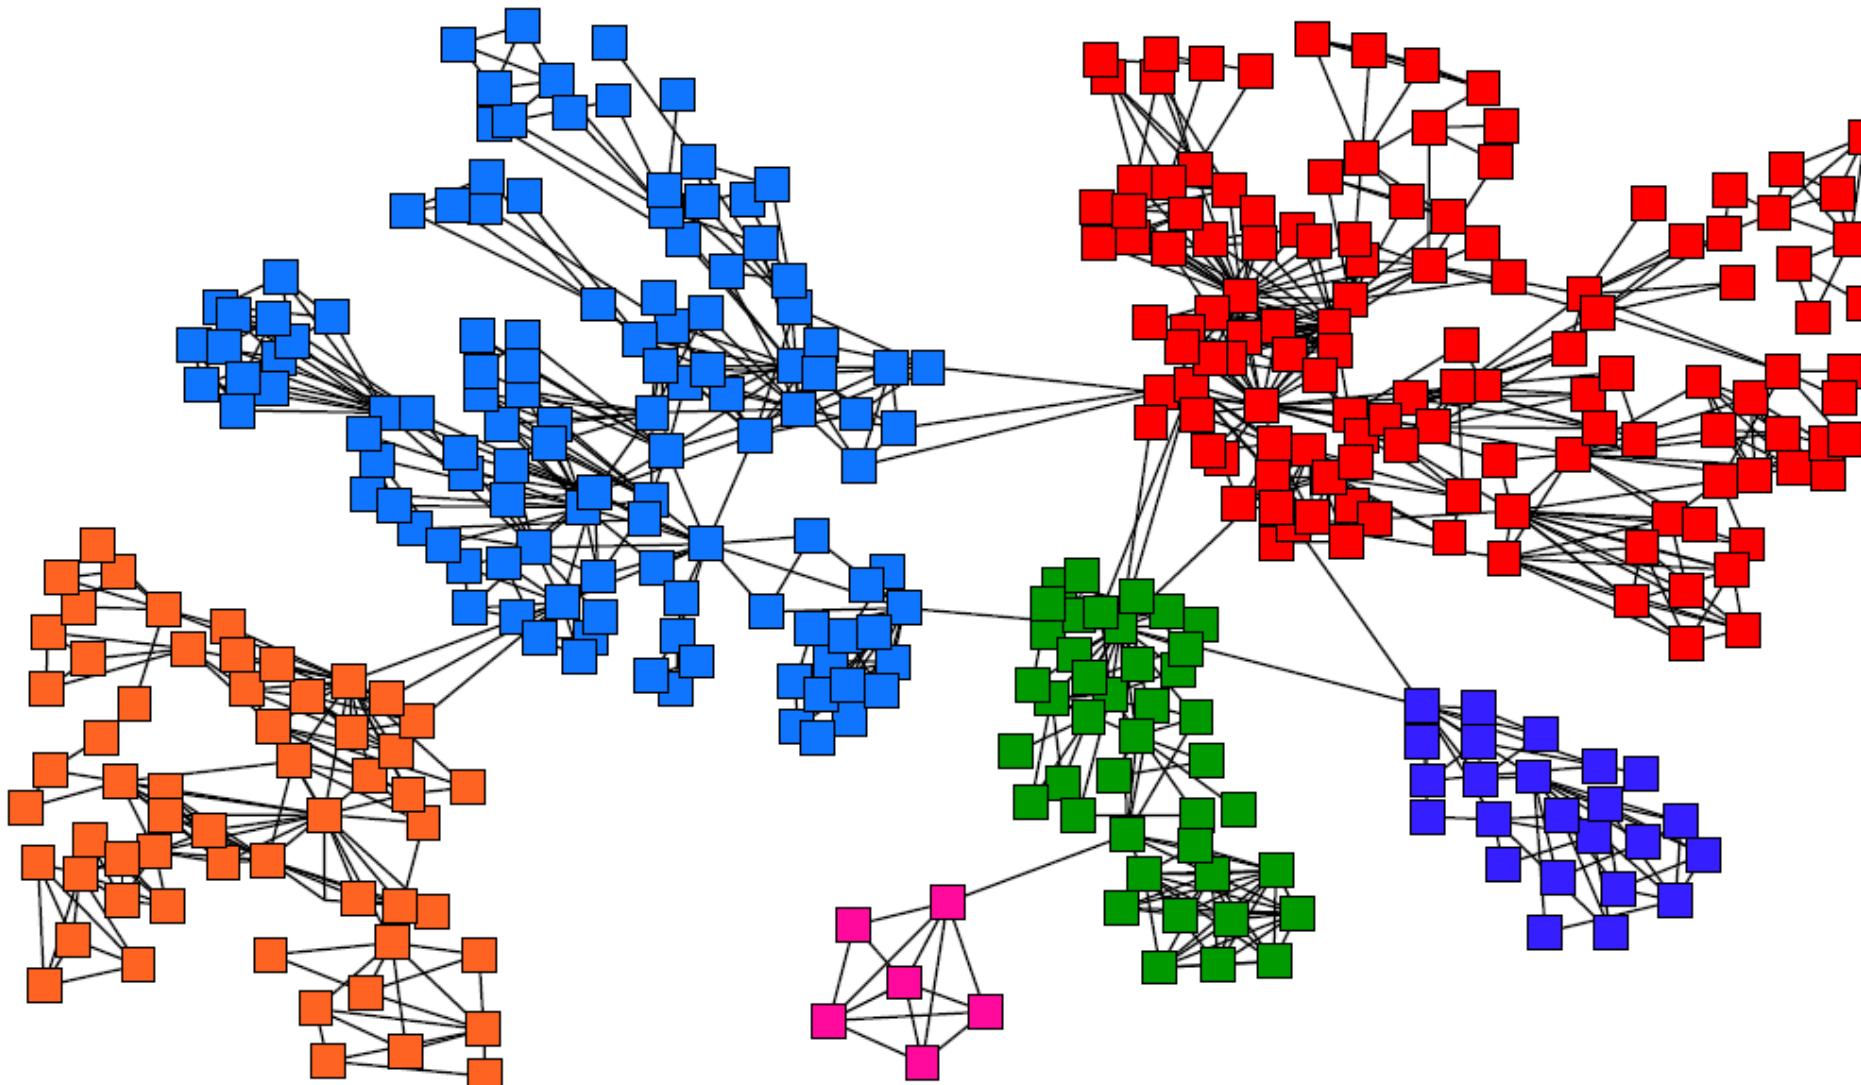
```
Social-Networks-Networkx.ipynb  
modularity_maximization.zip
```

Networks & Communities

- We often think of networks being organized into **modules, cluster, communities**:

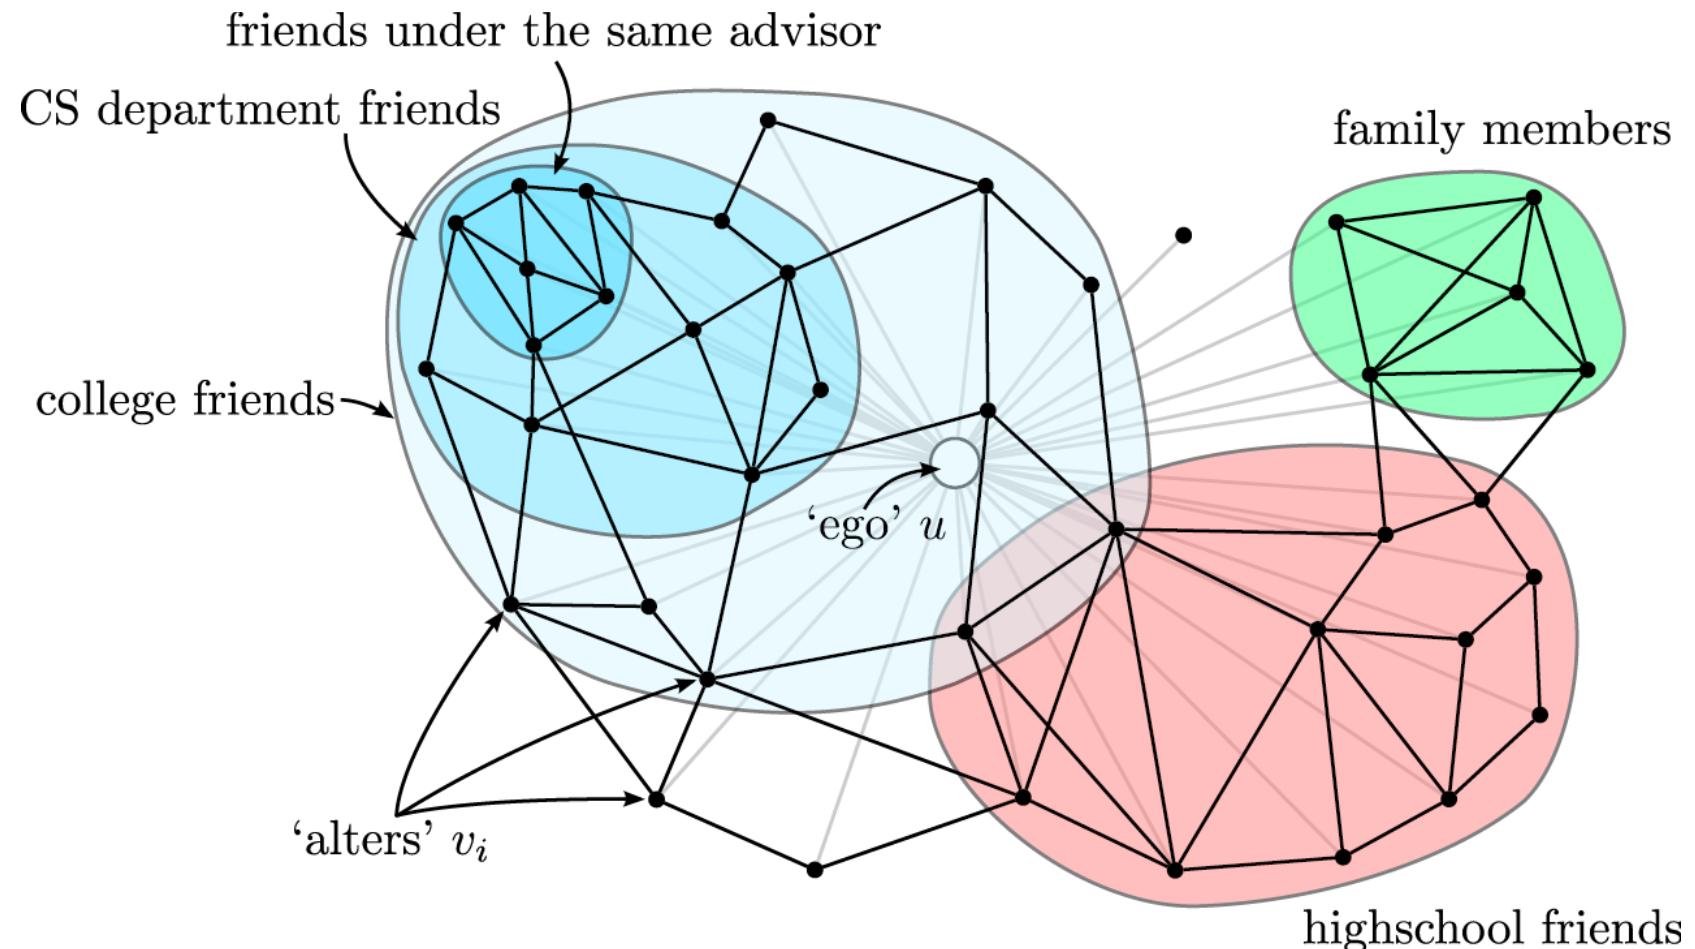


Goal: Find Densely Linked Clusters



Twitter & Facebook

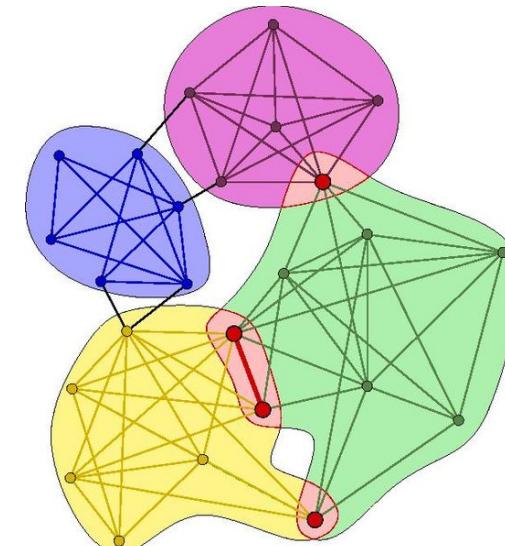
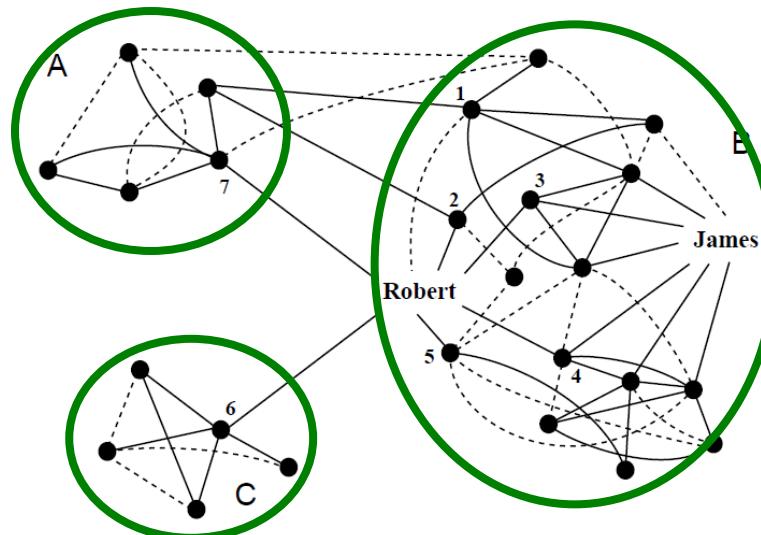
- **Discovering social circles, circles of trust:**



[McAuley, Leskovec: Discovering social circles in ego networks, 2012]

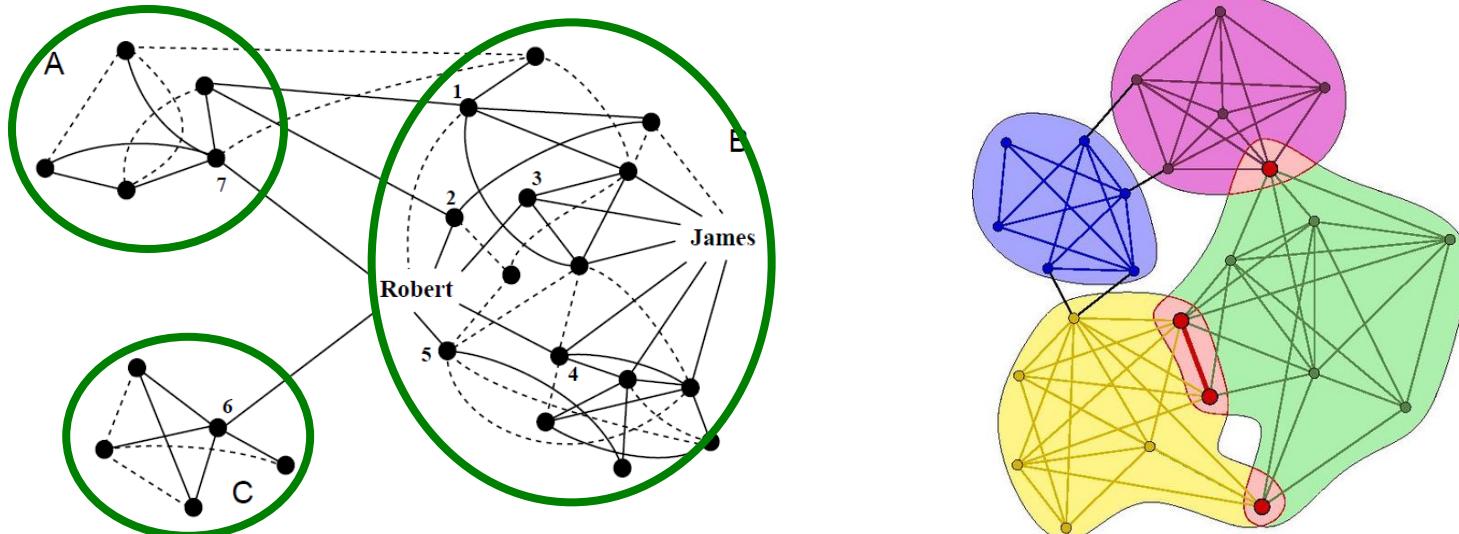
Community Detection (Graph Basics)

How to find communities?



Community Detection (Algorithms and Methods)

How to find communities?



We will work with **undirected** (unweighted) networks

Recall: Methods of Clustering

- **Hierarchical:**

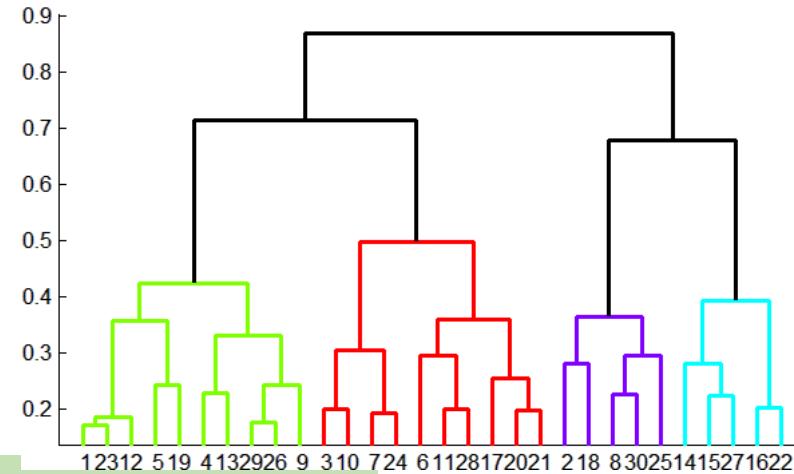
- **Agglomerative** (bottom up):

- Initially, each point is a cluster
 - Repeatedly combine the two “nearest” clusters into one
 - **Used a distance metric**

Today:

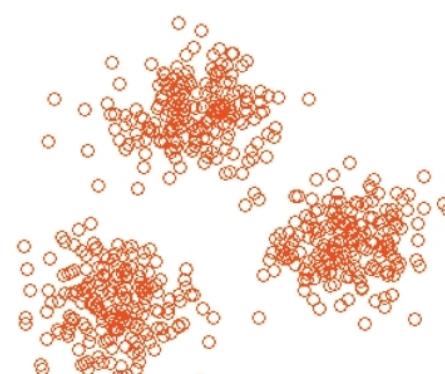
- **Divisive** (top down):

- Start with one cluster and recursively split it



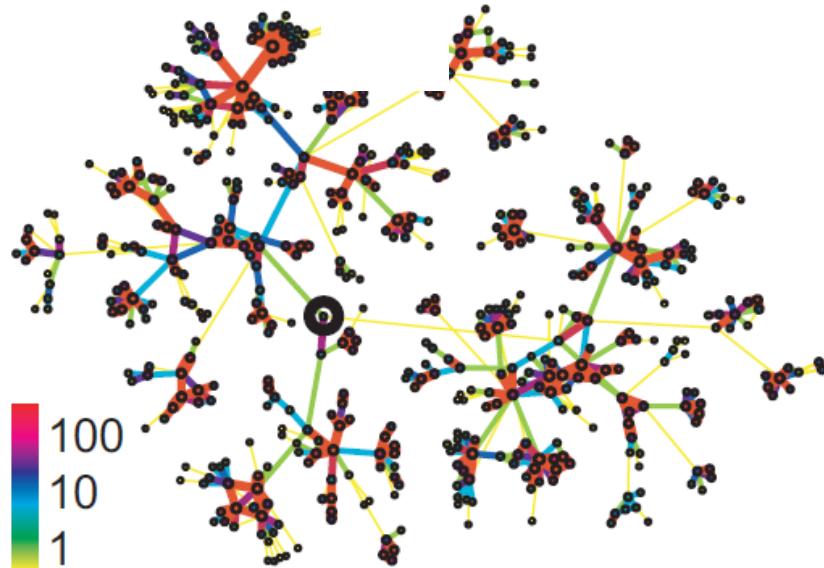
- **Point assignment:**

- Maintain a set of clusters
 - Points belong to “nearest” cluster
 - **Used a distance metric**

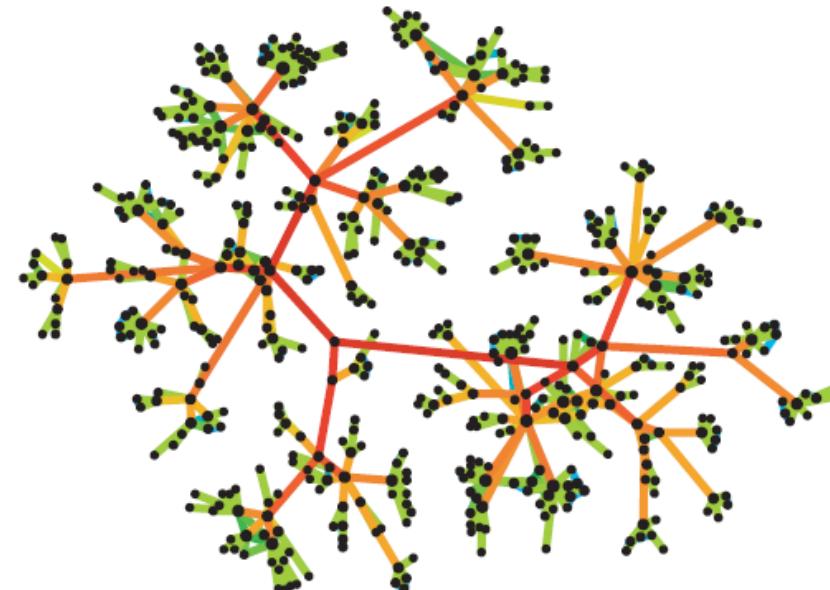


Betweenness Concept

- **Edge betweenness:** Number of shortest paths passing over the edge
- **Intuition:**



Edge strengths (call volume)
in a real network

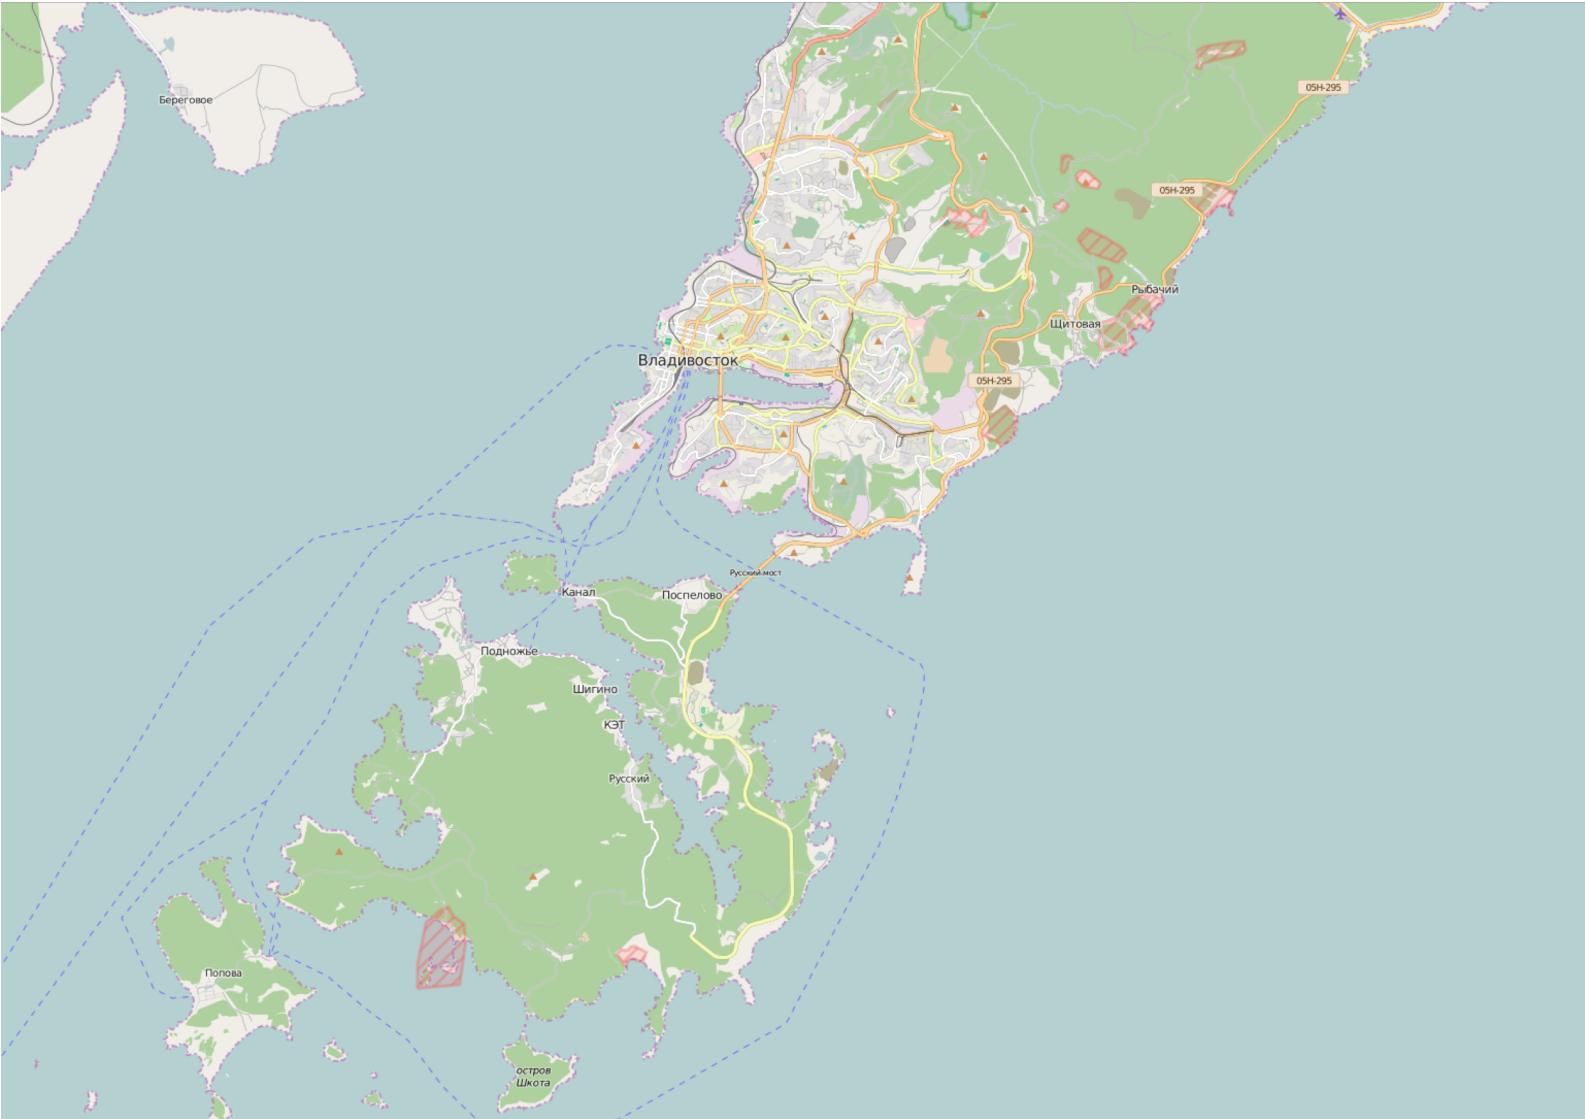


Edge betweenness
in a real network

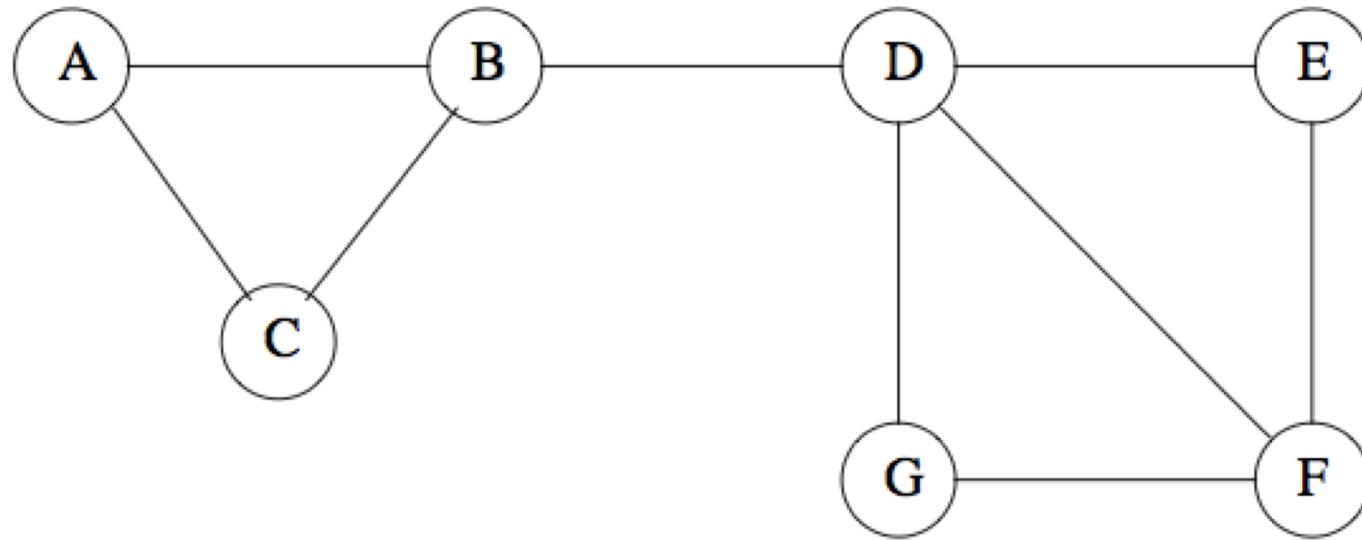
Betweenness Concept (Cont'd)

- Find edges in a social network graph that are least likely to be inside a community
- Betweenness of edge (a, b):
 - number of pairs of nodes x and y $\rightarrow x, y \in C$
 - **edge (a,b) lies on the shortest path between x and y**
- If there are several shortest paths between x and y, edge (a,b) is credited with **the fraction of those shortest paths** that include edge (a,b)
- **A high score is bad:** suggests that edge (a,b) runs between two different communities
 - **a and b are in different communities**

The Russian Bridge



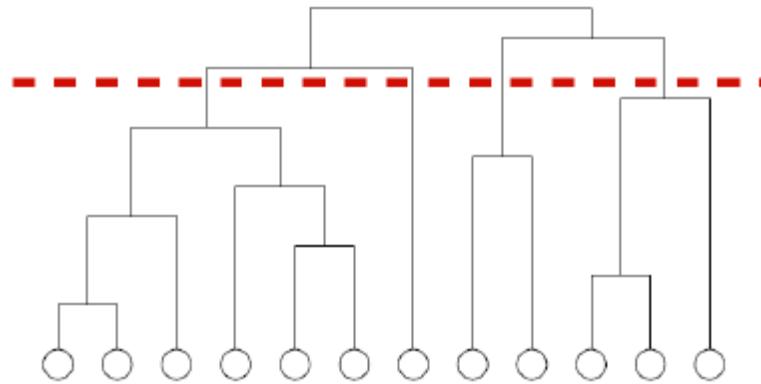
Betweenness Example



- Expect that edge (B,D) has highest betweenness
- (B,D) is on every shortest path from {A,B,C} to {D,E,F,G}
- Betweenness of (B,D) = $3 \times 4 = 12$
- (D,F) is on every shortest path from {A,B,C,D} to {F}
- Betweenness of (D,F) = $4 \times 1 = 4$
- **Natural communities: {A,B,C} and {D,E,F,G}**

We need to resolve 2
questions

1. **How to compute betweenness?**
2. How to select the number of clusters?



The Girvan-Newman Algorithm

- Want to **discover communities using divisive hierarchical clustering**
 - Start with one cluster (the social network) and recursively split it
- **Will do this** based on the notion of edge **betweenness**:
Number of shortest paths passing through the edge
- **Girvan-Newman Algorithm:**
 - Visits each node X once
 - Computes the number of shortest paths from X to **each of the other nodes** that go through each of the edges
- **Repeat:**
 - Calculate betweenness of edges
 1. Thresholding to remove high betweenness edges, or
 2. Remove edges with highest betweenness: **between** communities
- **Connected components are communities**
- Gives a hierarchical decomposition of the network

Girvan-Newman Algorithm (1)

- Visit each node X once and **compute the number of shortest paths from X to each of the other nodes that go through each of the edges**
- **1) Perform a breadth-first search (BFS) of the graph, starting at node X**
 - The level of each node in BFS is length of the shortest path from X to that node
 - So edges that go between nodes on the same level can never be part of a shortest path from X
 - **Edges between levels are called DAG edges** (DAG = Directed Acyclic Graph)
 - **Each DAG edge is part of at least one shortest path from root X**

Girvan-Newman Algorithm (2)

- 2) Label each node by the number of shortest paths that reach it from the root node
 - Example: BFS starting from node E, labels assigned

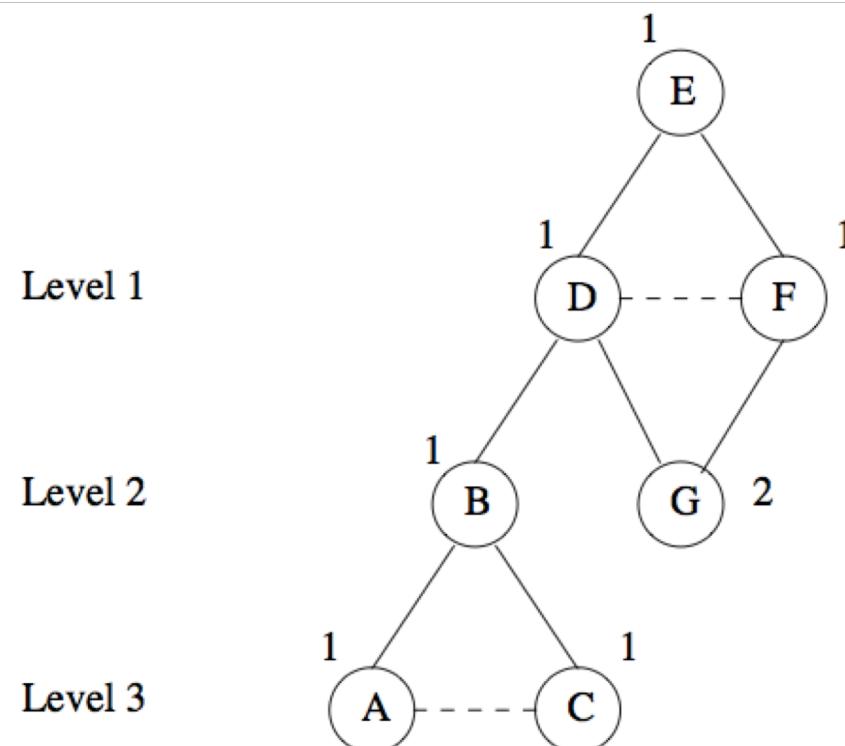
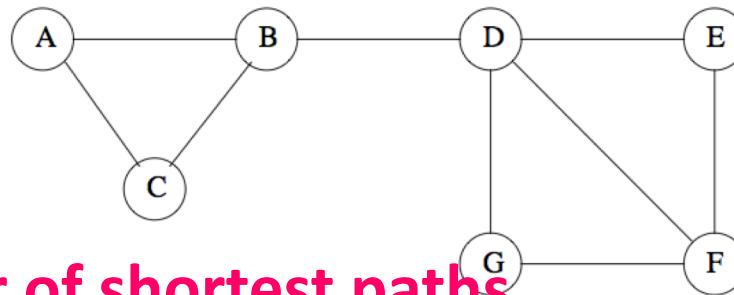


Figure 10.4: Step 1 of the Girvan-Newman Algorithm

Girvan-Newman Algorithm (3)

- 3) Calculate for each edge e , the sum over all nodes Y (of the fraction) of the shortest paths from the root X to Y that go through edge e
 - Compute this sum for nodes and edges, starting from the bottom of the graph
 - Each **node** other than the root node is given a credit of 1
 - Each **leaf node** in the DAG gets a credit of 1
 - Each **node** that is not a leaf gets credit = 1 + sum of credits of the DAG edges from that node to level below
 - A DAG **edge** e entering node Z (from the level above) is given a share of the credit of Z proportional to the fraction of shortest paths from the root to Z that go through e

Girvan-Newman Algorithm (4)

- Assign node and edge values starting from bottom

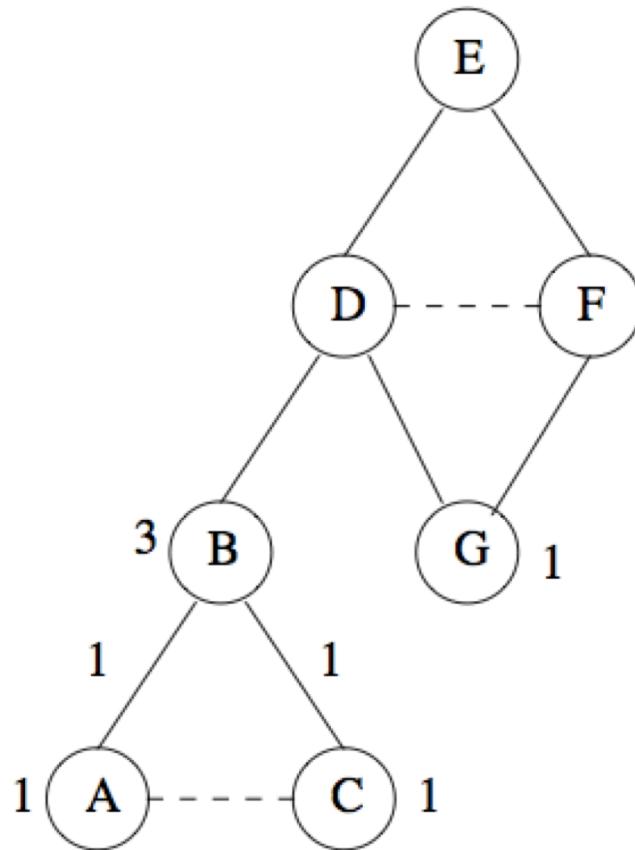
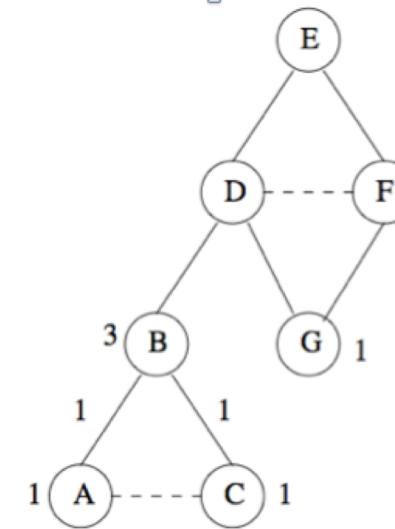


Figure 10.5: Final step of the Girvan-Newman Algorithm – levels 3 and 2

Girvan-Newman Algorithm (5)

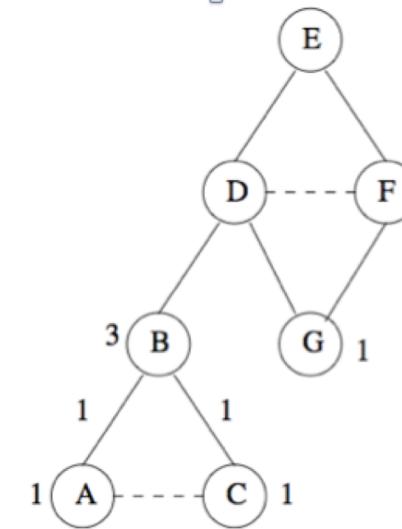
Assigning credits:

- A and C are **leaves**: get credit = 1
- Each of these nodes **has only one parent**, so their credit=1 is **given to edges** (B,A) and (B,C)
- At level 2, G is a **leaf**: gets credit = 1
- B gets credit **1 + credit of DAG edges entering from below**
 $= 1 + 1 + 1 = 3$
- B has only **one parent**, so edge (D,B) **gets entire credit of node B = 3**
- **Node G has 2 parents (D and F): how do we divide credit of G between the edges?**



Girvan-Newman Algorithm (6)

- In this case, **both D and F have just one path from E to each of those nodes**
 - So, **give half credit of node G to each of those edges**
 - Credit = $1/(1 + 1) = 0.5$
- **In general, how we distribute credit of a node to its edges depends on number of shortest paths**
 - Say there were 5 shortest paths to D and only 3 to F
 - Then credit of edge (D,G) = $5/8$ and credit of edge (F,G) = $3/8$
- Node D gets credit = **1 + credits of edges below it** =
 $1 + 3 + 0.5 = 4.5$
- Node F gets credit = $1 + 0.5 = 1.5$
- D has **only one parent**, so Edge (E,D) gets credit = 4.5 from D
- Likewise for F: Edge (E,F) gets credit = 1.5 from F



Girvan-Newman Algorithm (7): Completion of Credit Calculation starting at node E

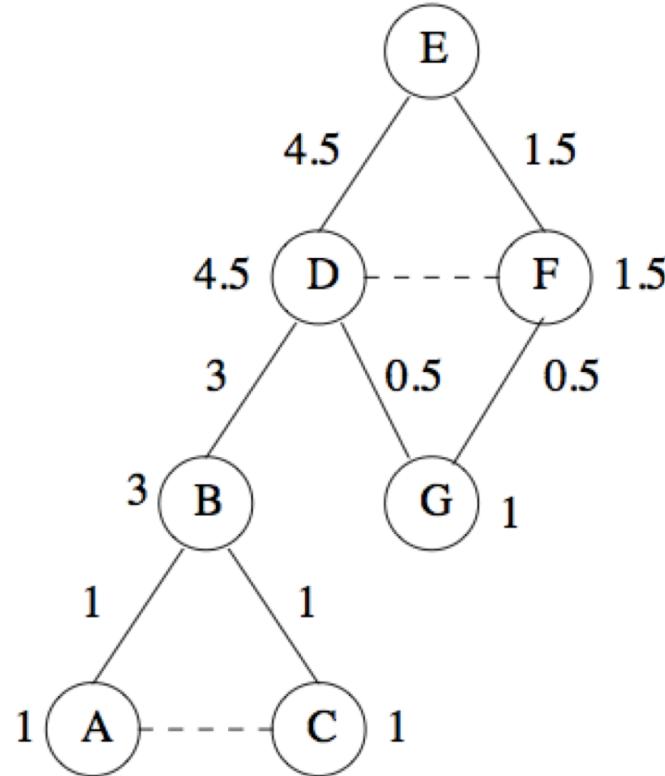


Figure 10.6: Final step of the Girvan-Newman Algorithm – completing the credit calculation

Girvan-Newman Algorithm (8): Overall Betweenness Calculation

- To complete betweenness calculation, must:
 - Repeat this for every node as root
 - Sum the contributions on each edge
 - Divide by 2 to get true betweenness
 - since every shortest path will be counted twice, once for each of its endpoints

Using Betweenness to Find Communities: Clustering

- Betweenness scores for edges of a graph behave something like a distance metric
 - Not a true distance metric
- Could cluster by taking edges in increasing order of betweenness and adding to graph one at a time
 - At each step, connected components of graph form clusters
- Girvan-Newman: Start with the graph and all its edges and remove edges with highest betweenness
 - Continue until graph has broken into suitable number of connected components
 - Divisive hierarchical clustering (top down)
 - Start with one cluster (the social network) and recursively split it

Using Betweenness to Find Communities (2)

- (B,D) has highest betweenness (12)
- Removing edge would give natural communities we identified earlier: {A,B,C} and {D,E,F,G}

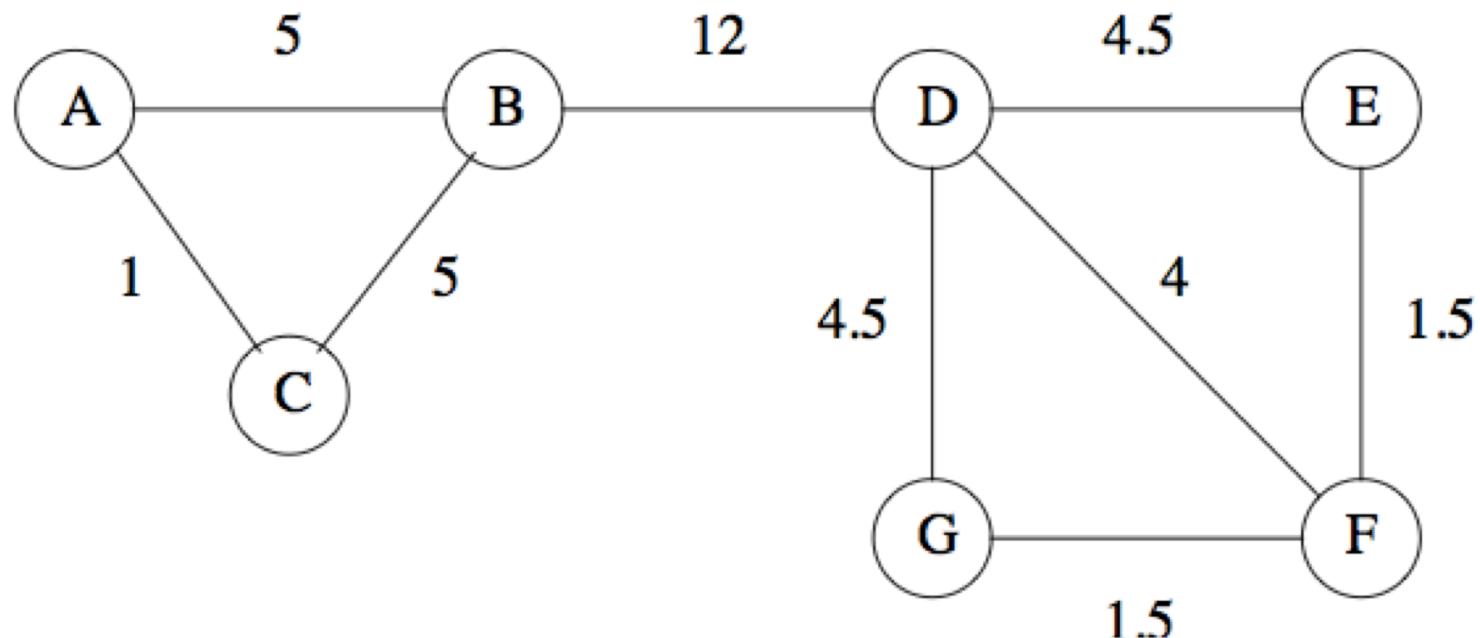


Figure 10.7: Betweenness scores for the graph of Fig. 10.1

Using Betweenness to Find Communities (3): Thresholding

- Could continue to remove edges with highest betweenness

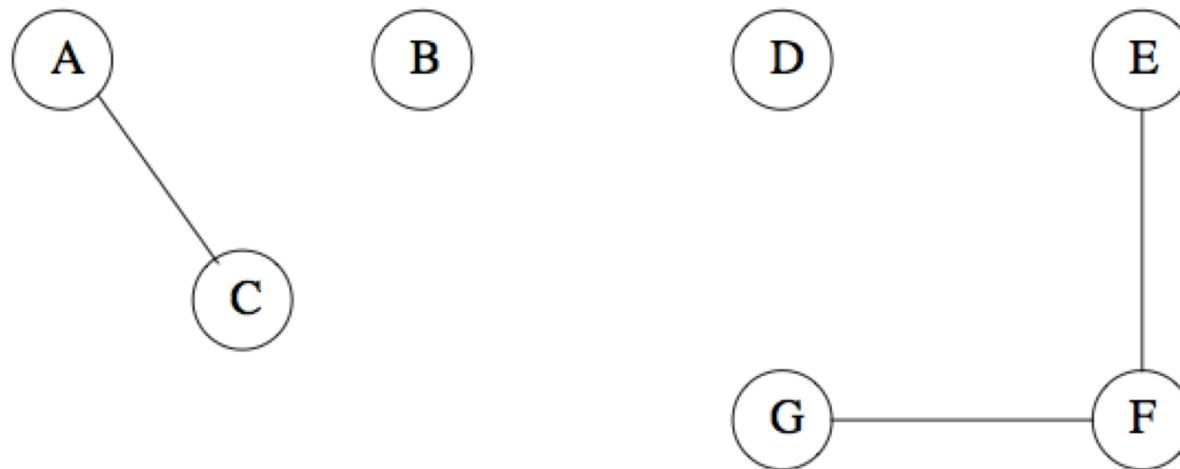
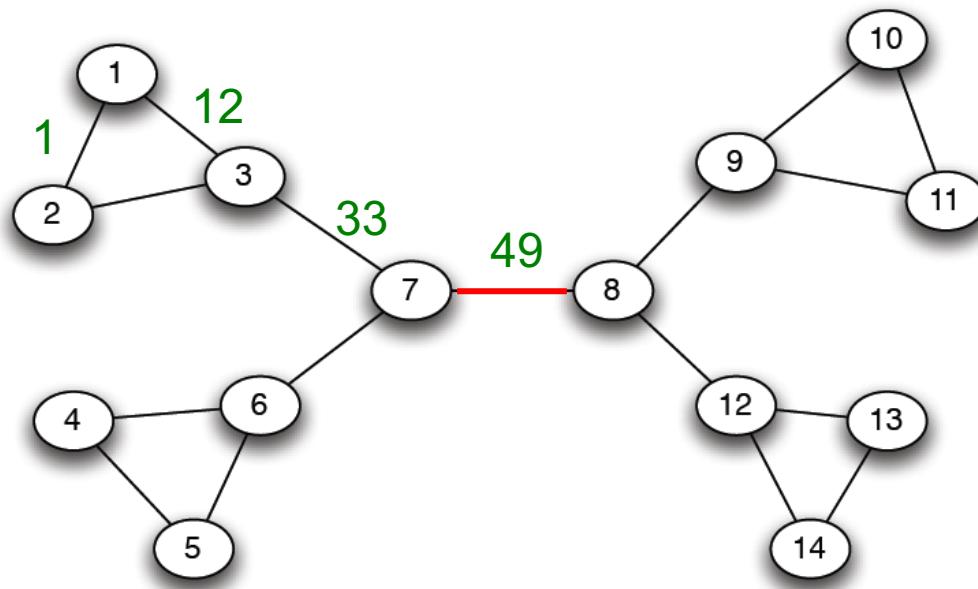


Figure 10.8: All the edges with betweenness 4 or more have been removed

Run Girvan-Newman Iteratively for Community Detection

- Recall: Divisive hierarchical clustering based on the notion of edge **betweenness**:
Number of shortest paths passing through the edge
- **Girvan-Newman Algorithm:**
 - Undirected unweighted networks
 - **Repeat until no edges are left:**
 - Calculate betweenness of edges
 - **This time: remove edges with highest betweenness**
 - Connected components are communities
 - Gives a hierarchical decomposition of the network

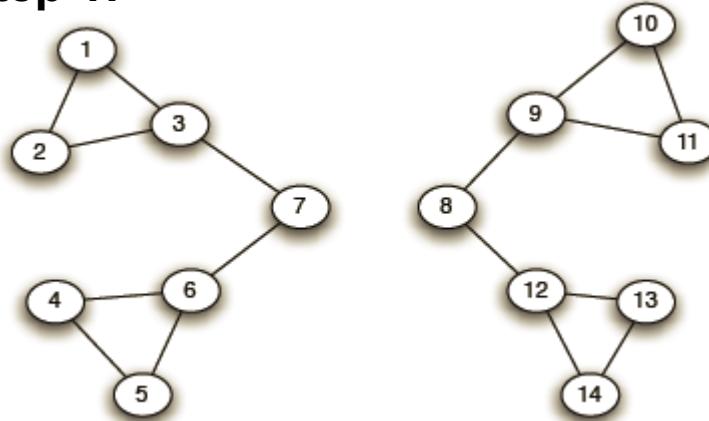
Girvan-Newman: Example



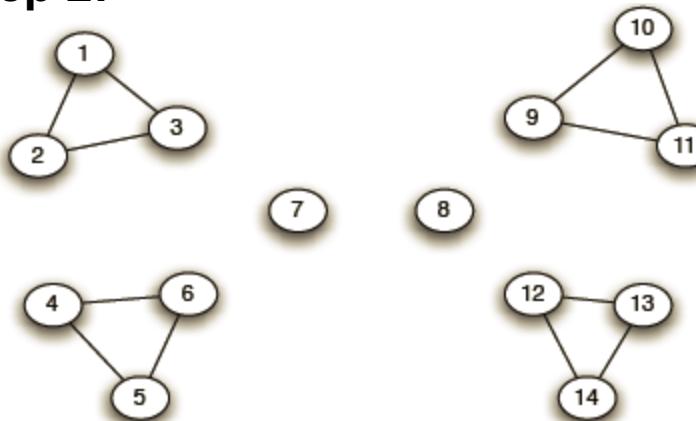
Need to re-compute
betweenness at
every step

Girvan-Newman: Example

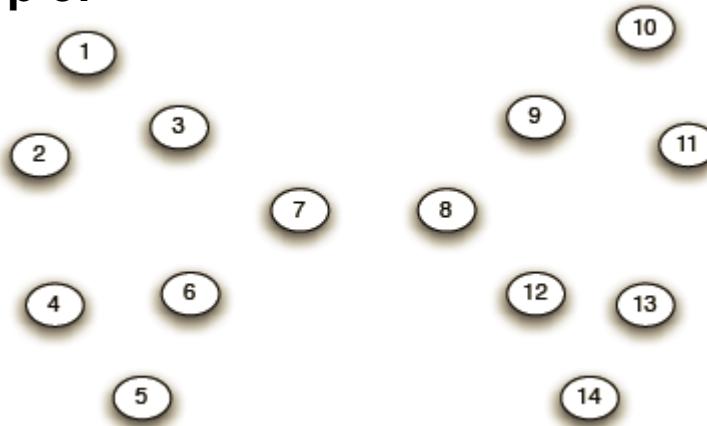
Step 1:



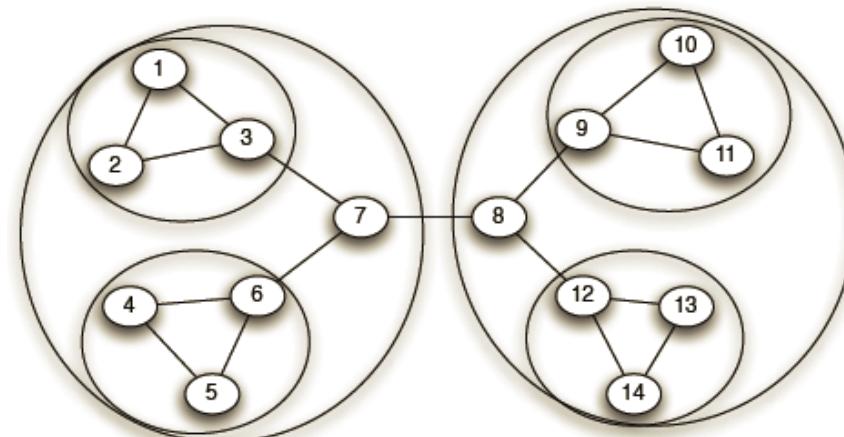
Step 2:



Step 3:

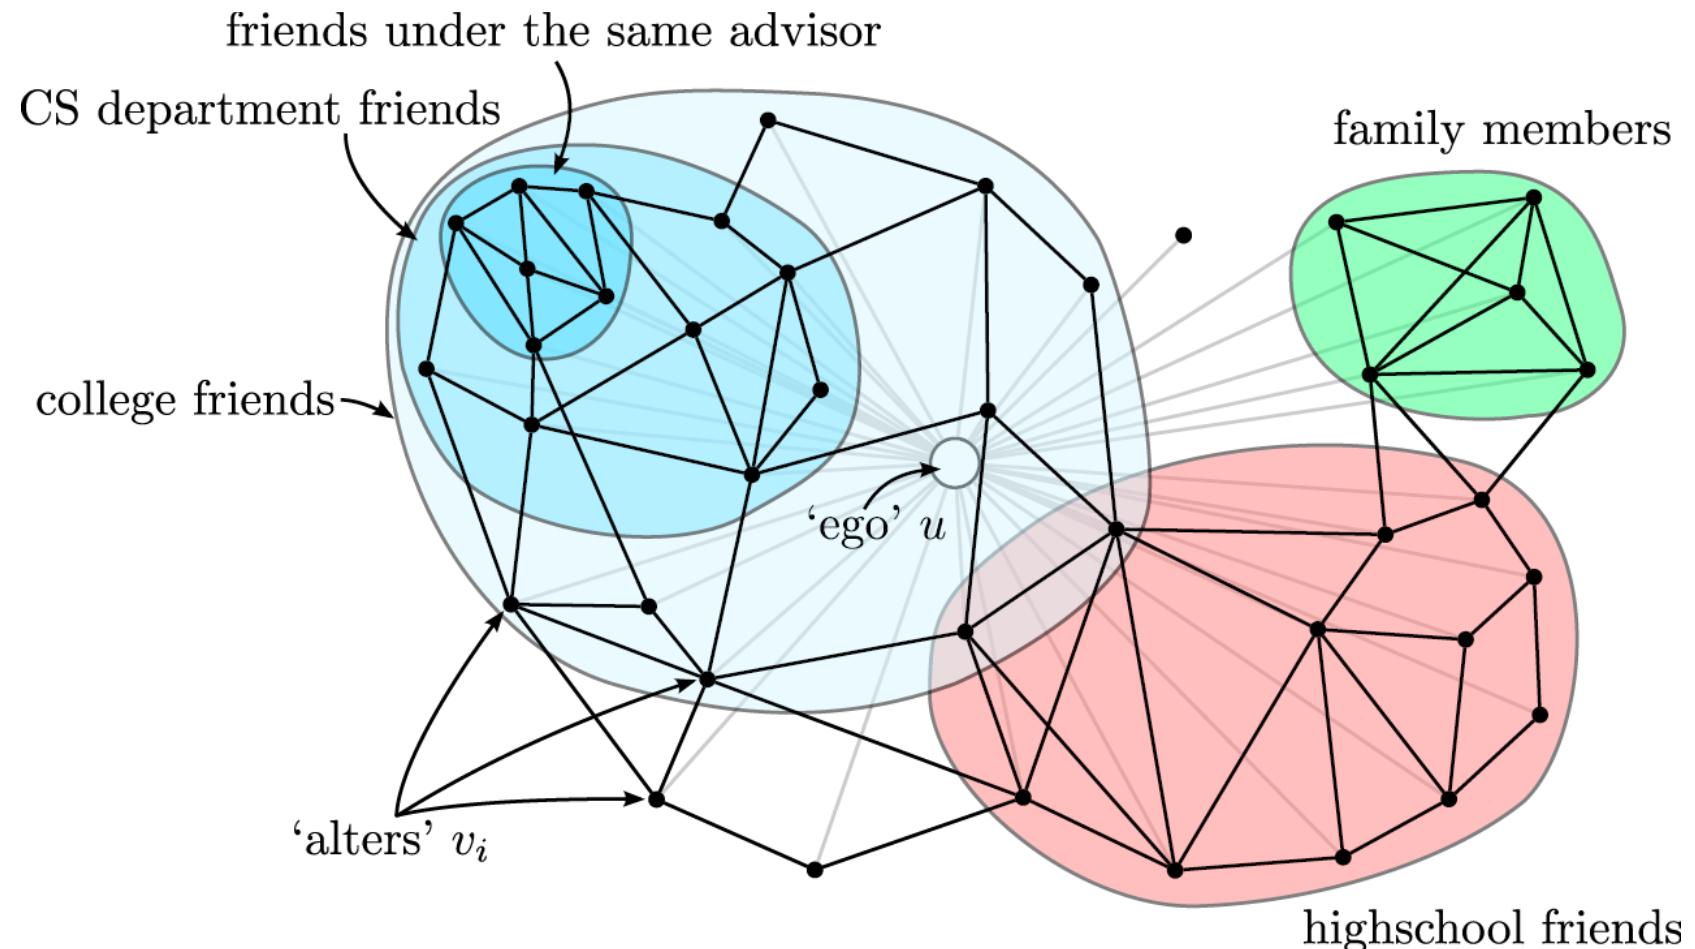


Hierarchical network decomposition:



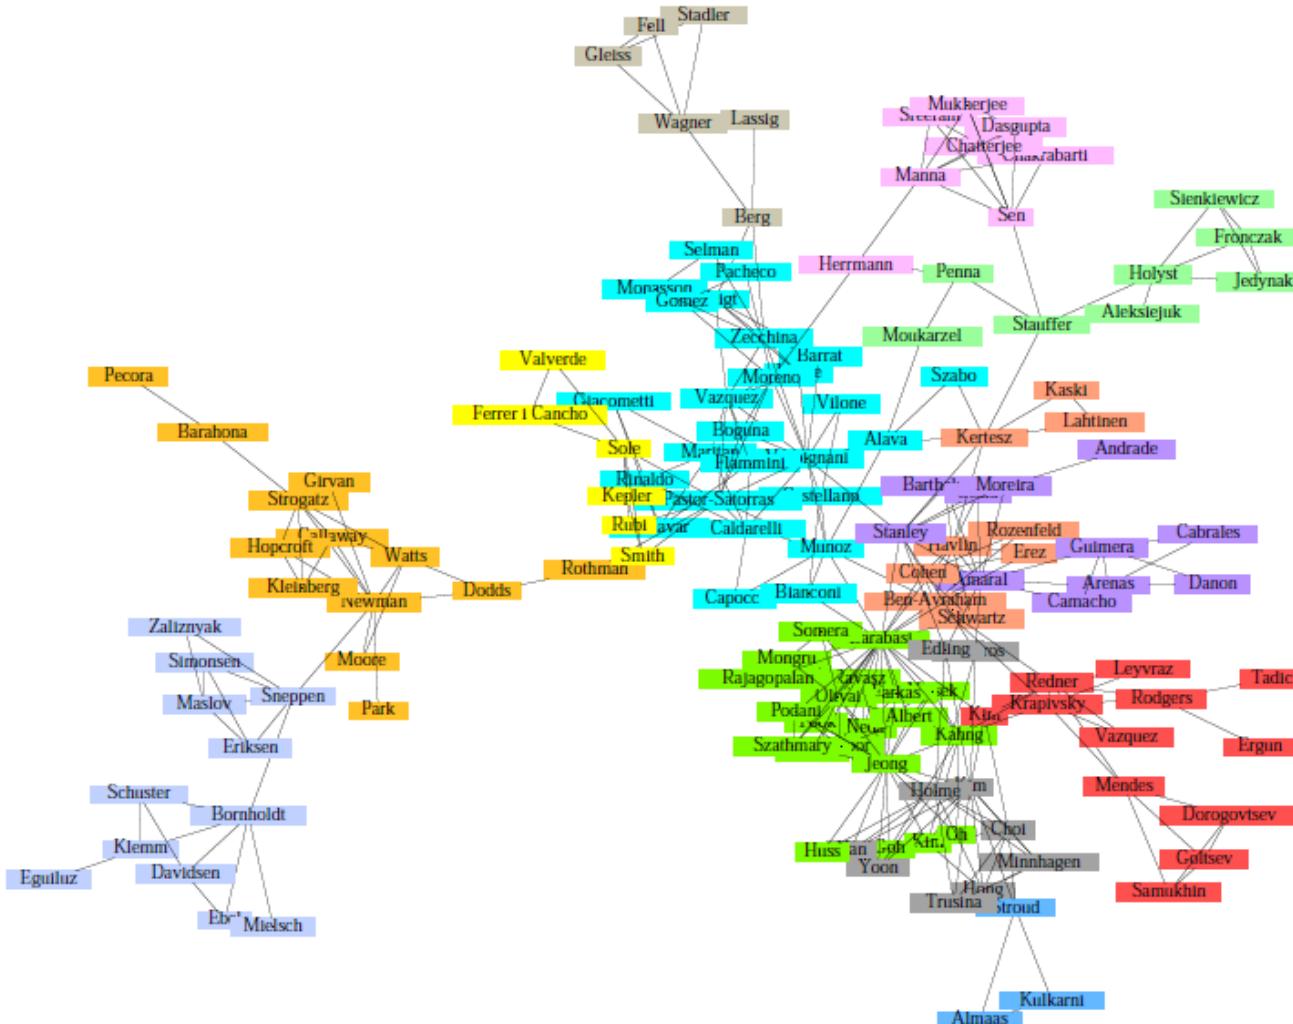
Recall: Twitter & Facebook

- **Discovering social circles, circles of trust:**



[McAuley, Leskovec: Discovering social circles in ego networks, 2012]

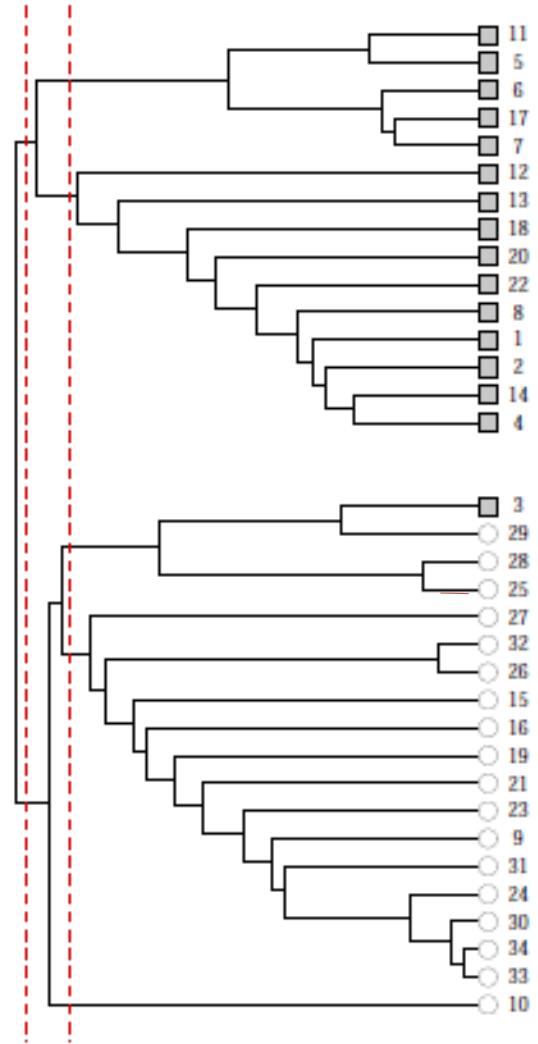
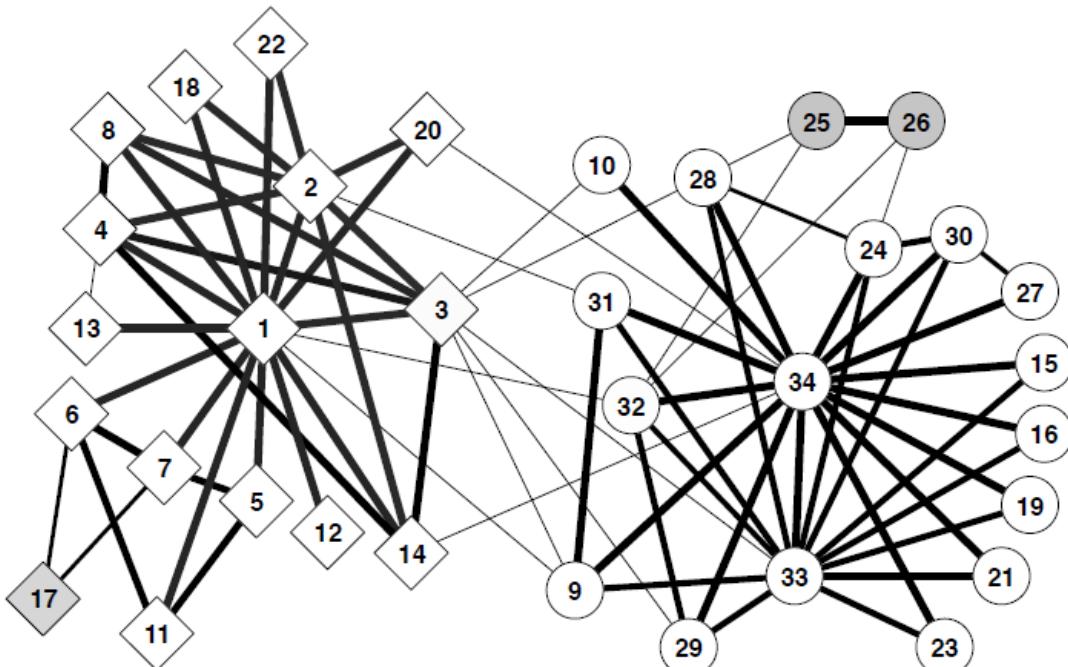
Girvan-Newman: Results



Communities in physics collaborations

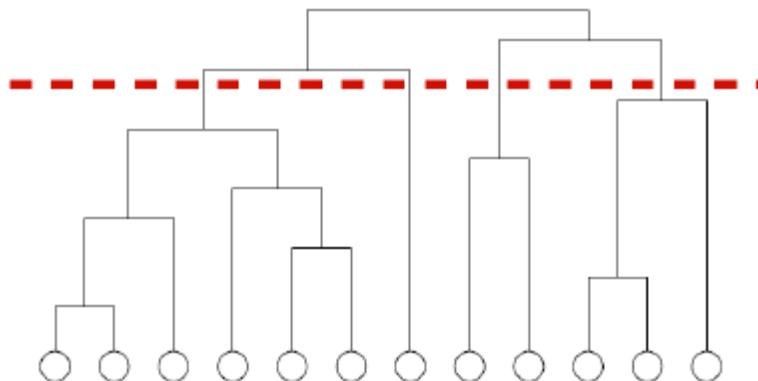
Girvan-Newman: Results

- Zachary's Karate club:
Hierarchical decomposition



We need to resolve 2 questions

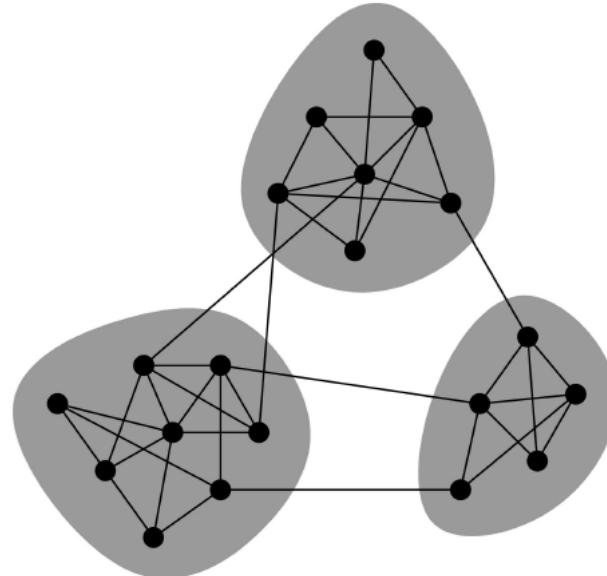
1. How to compute betweenness?
2. How to select the number of clusters?



Network Communities

- **Communities:** sets of tightly connected nodes
- Define: **Modularity Q**
 - A measure of how well a network is partitioned into communities
 - Given a partitioning of the network into groups $\in_{\epsilon} S$:

$$Q = \sum_{s \in S} [(\# \text{ edges within group } s) - (\text{expected } \# \text{ edges within group } s)]$$



Need a null model!

The null model is a graph which matches one specific graph in some of its structural features, but which is otherwise taken to be an instance of a random graph. The null model is used as a term of comparison, to verify whether the graph in question displays some feature, such as community structure, or not.

Modularity

- Modularity of partitioning S of graph G :

- $Q = \sum_{s \in S} [(\# \text{ edges within group } s) - (\text{expected } \# \text{ edges within group } s)]$

- $Q(G, S) = \frac{1}{2m} \sum_{s \in S} \sum_{i \in s} \sum_{j \in s} \left(A_{ij} - \frac{k_i k_j}{2m} \right)$

Normalizing cost.: $-1 < Q < 1$

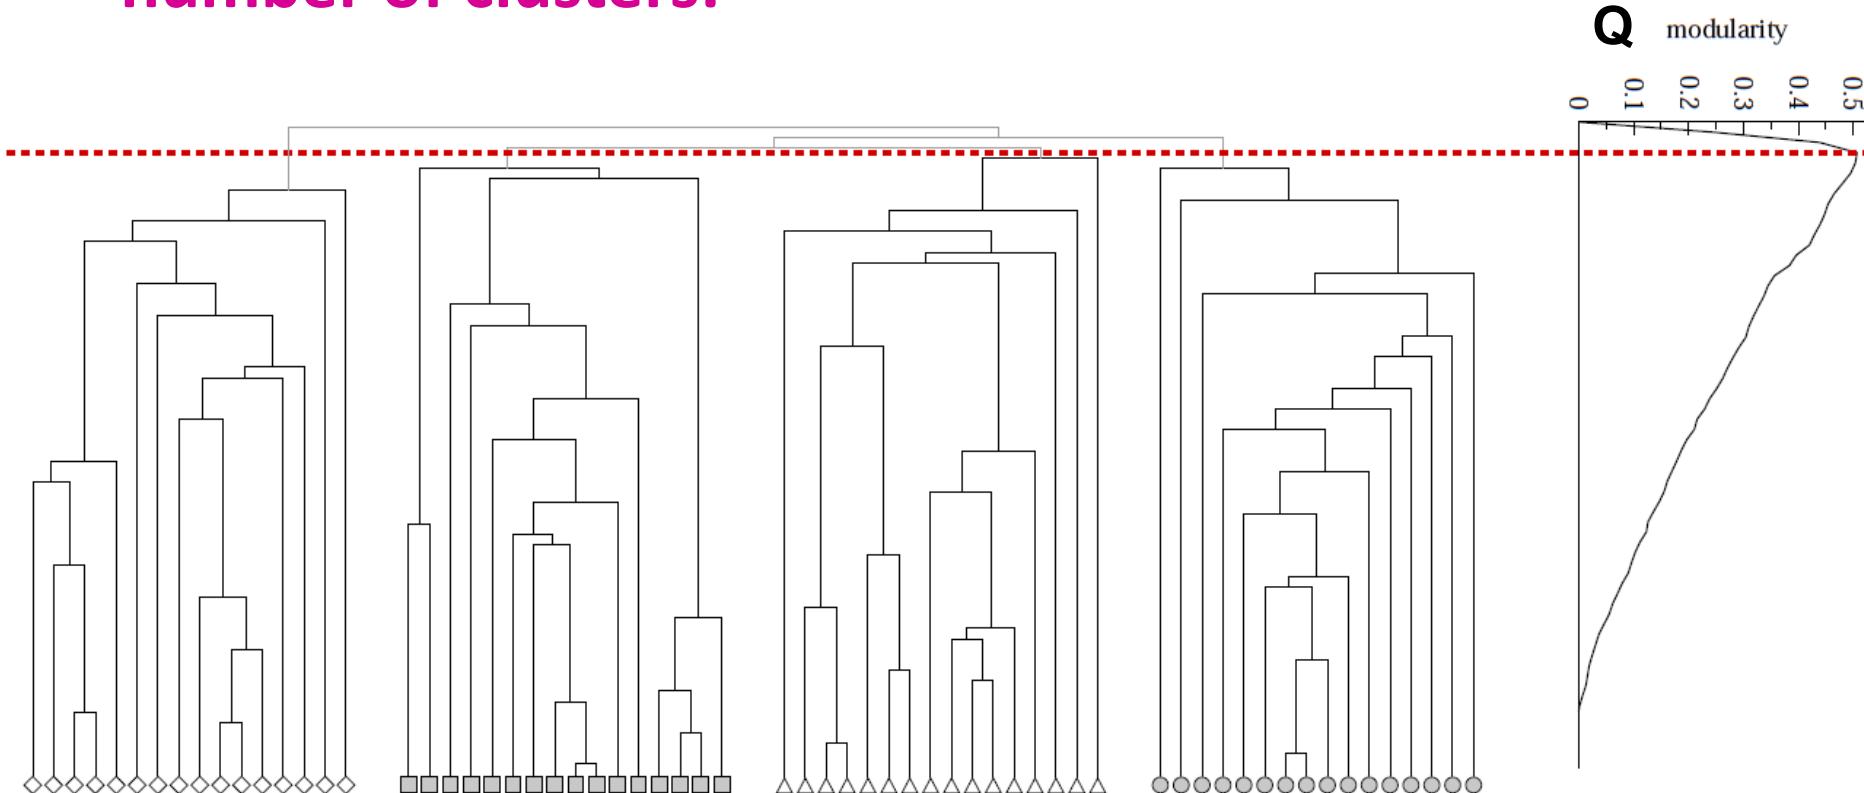
$A_{ij} = 1$ if i connects j ,
0 else

- Modularity values take range $[-1, 1]$

- It is positive if the number of edges within groups exceeds the expected number
- $0.3 - 0.7 < Q$ means significant community structure

Modularity: Number of clusters

- Modularity is useful for selecting the number of clusters:



Spectral Clustering

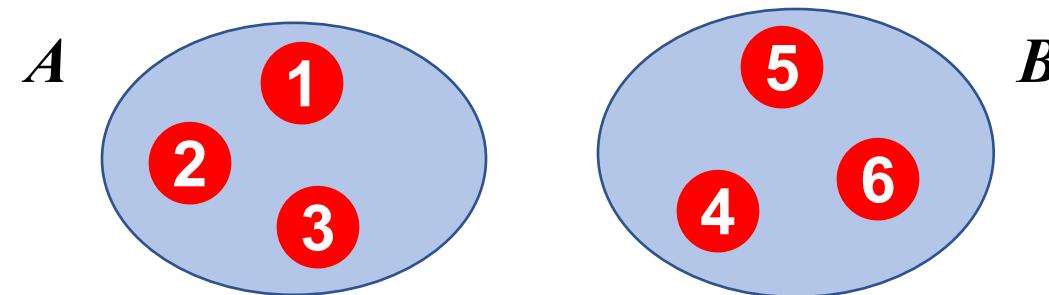
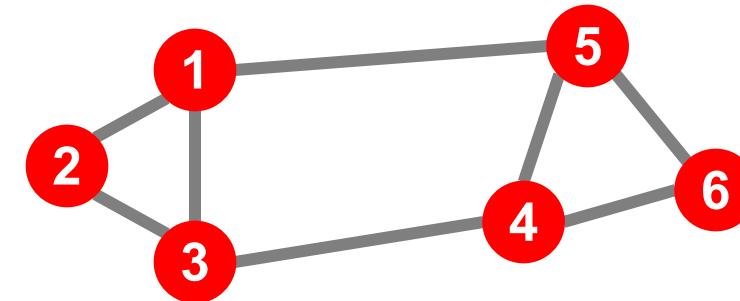
Another approach to organizing social-network graphs

Partitioning Graphs

- Another approach to organizing social networking graphs
- **Problem:** partitioning a graph to minimize the number of edges that connect different components (communities)
- Goal of minimizing the cut size
- If you just joined Facebook with only one friend
 - Don't want to partition the graph with you disconnected from rest of the world
 - Want components to be not too unequal in size

Graph Partitioning

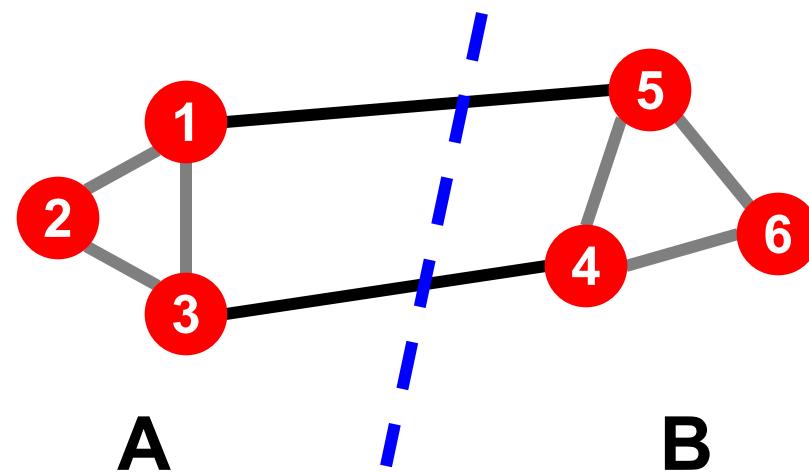
- Undirected graph
- Bi-partitioning task:
 - Divide vertices into two disjoint groups



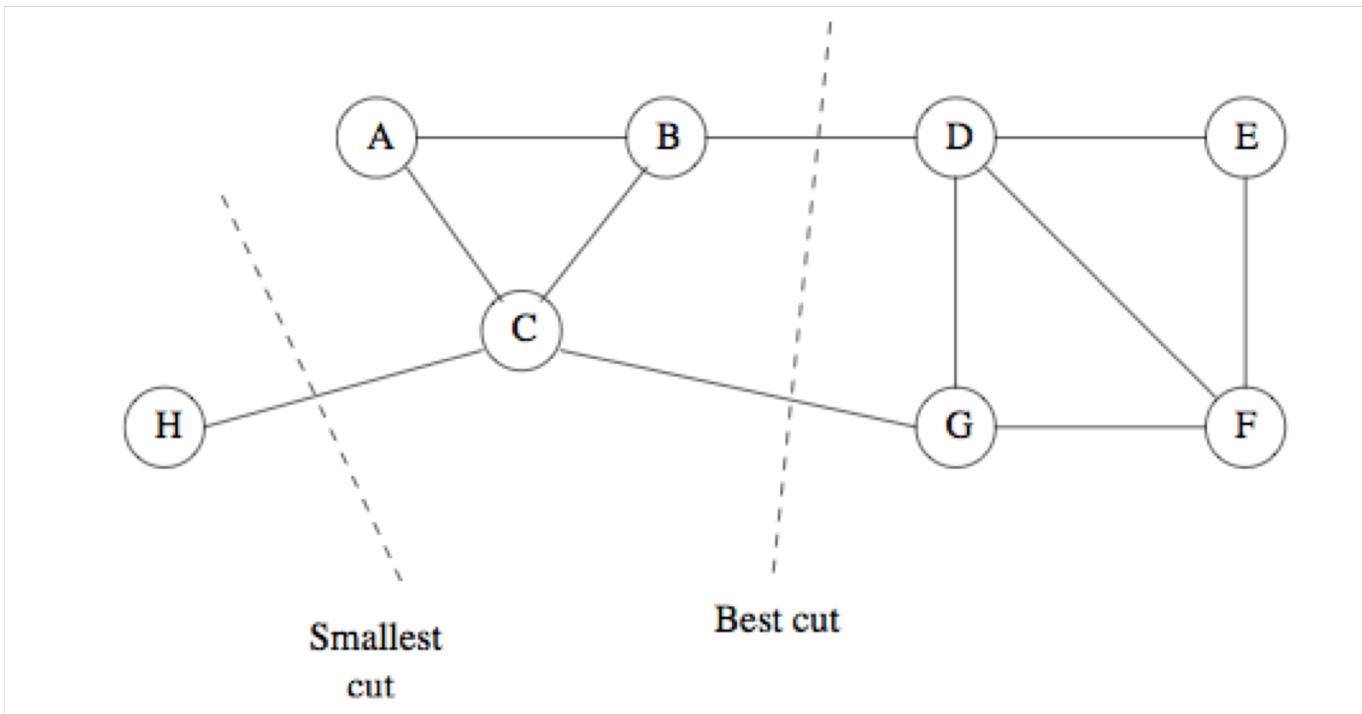
- Questions:
 - How can we define a “good” partition of ?
 - How can we efficiently identify such a partition?

Graph Partitioning

- What makes a good partition?
 - Divide nodes into two sets so that the **cut (set of edges that connect nodes in different sets) is minimized**
 - Want the two sets to be approximately equal in size
 - Maximize the number of within-group connections
 - Minimize the number of between-group connections



Example 10.14

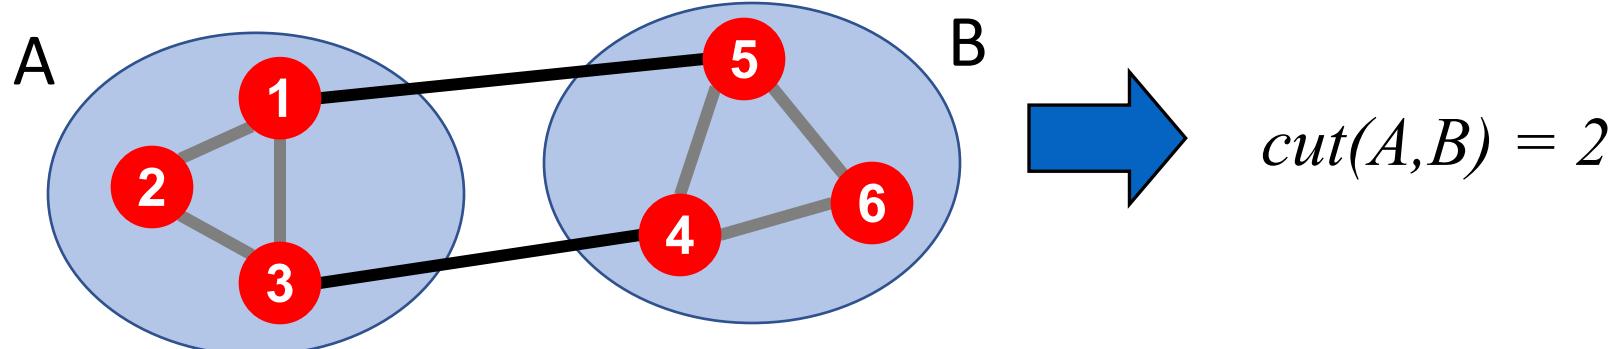


- If we minimize cut: best choice is to put H in one set, other nodes in other set
- But: we reject partitions where one set is too small
- Better is to use cut with (B,D) and (C,G)
- **Smallest cut is not necessarily the best cut**

Graph Cuts

- Express partitioning objectives as a function of the “edge cut” of the partition
- **Cut:** Set of edges with only one vertex in a group:

$$cut(A, B) = \sum_{i \in A, j \in B} w_{ij}$$



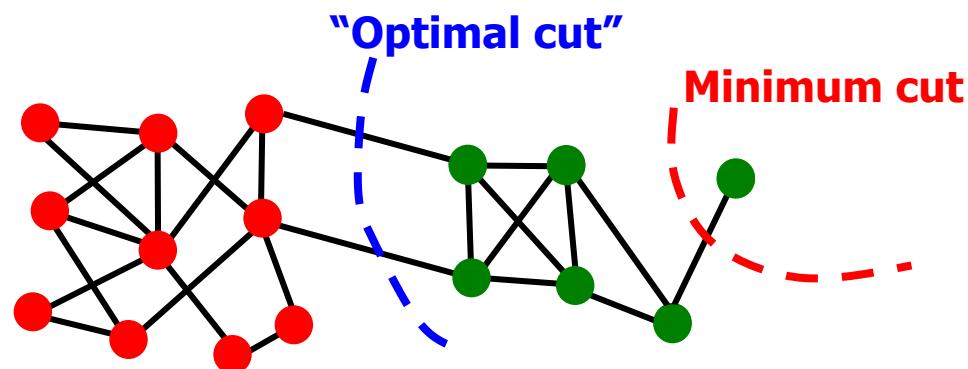
Graph Cut Criterion

- Criterion: **Minimum-cut**
 - Minimize weight of connections between groups

- Degenerate case:

$$\arg \min_{A,B} \text{cut}(A,B)$$

- Problem:
 - Only considers external cluster connections
 - Does not consider internal cluster connectivity



Graph Cut Criteria

- Criterion: **Normalized-cut** [Shi-Malik, '97]
 - Connectivity between groups relative to the density of each group

$$ncut(A, B) = \frac{cut(A, B)}{vol(A)} + \frac{cut(A, B)}{vol(B)}$$

$vol(A)$: total weight of the edges with at least

one endpoint in A : $vol(A) = \sum_{i \in A} k_i$: $vol(A) = \sum_{i \in A} k_i$

■ Why use this criterion?

- Produces more balanced partitions

• How do we efficiently find a good partition?

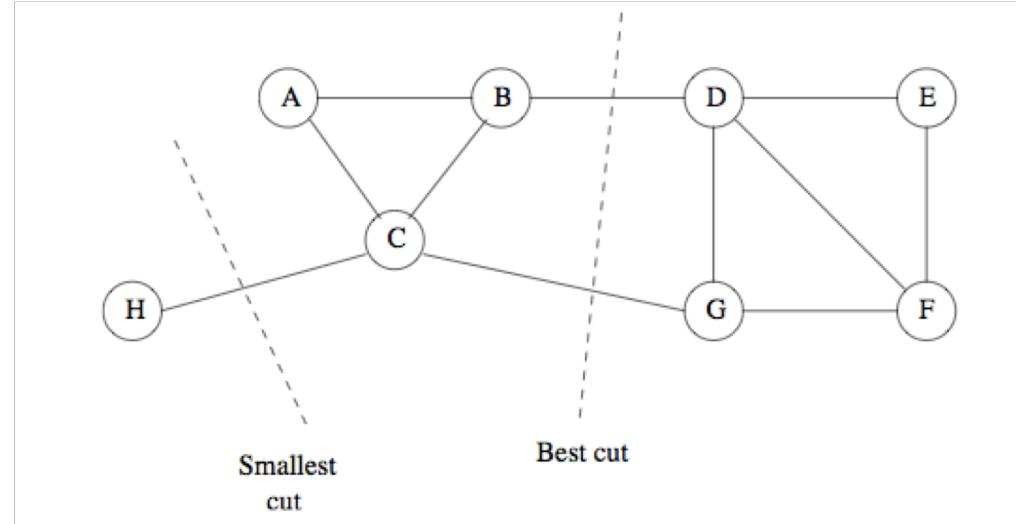
- Problem: Computing optimal cut is NP-hard

Example 10.15

- Partition nodes of graph into two disjoint sets S and T
- Normalized Cut for S and T is:

$$\frac{\text{Cut}(S,T) + \text{Cut}(S,T)}{\text{Vol}(S) \quad \text{Vol}(T)}$$

- If we choose $S=\{H\}$ and $T=\{A,B,C,D,E,F,G\}$ then $\text{Cut}(S,T) = 1$
 - $\text{Vol}(S) = 1$ (number of edges with at least one end in S)
 - $\text{Vol}(T) = 11$: all edges have at least one node in T
 - Normalized cut is $1/1 + 1/11 = 1.09$**
- For cut (B,D) and (C,G): $S = \{A,B,C,H\}$, $T = \{D,E,F,G\}$, $\text{Cut}(S,T) = 2$
- $\text{Vol}(S) = 6$, $\text{Vol}(T) = 7$, **normalized cut: $2/6 + 2/7 = 0.62$**



Using Matrix Algebra to Find Good Graph Partitions

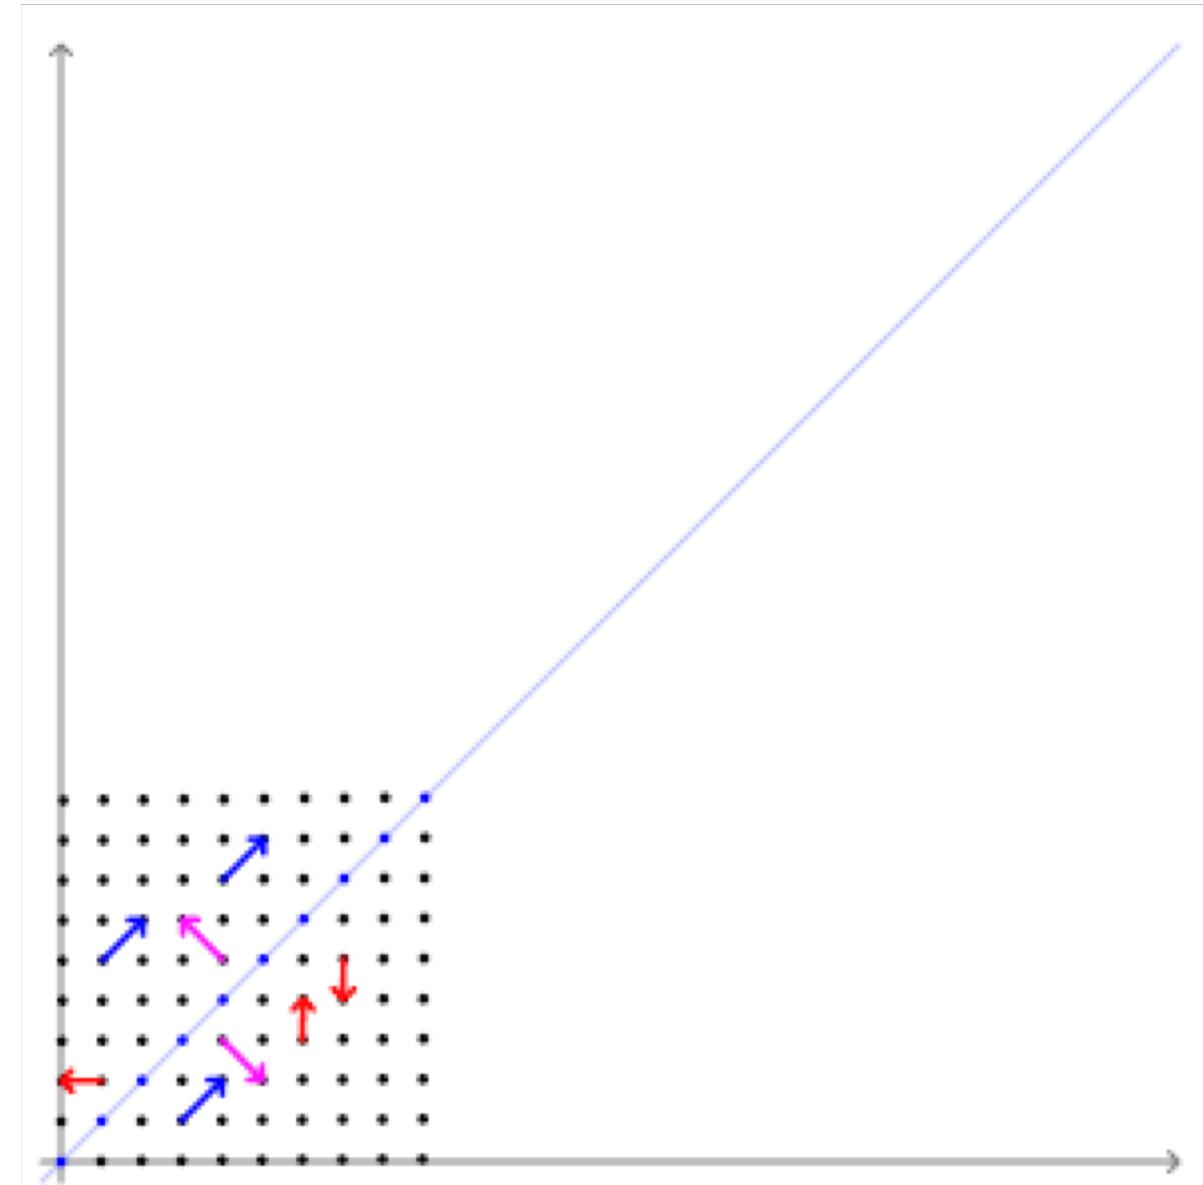
- Three matrices that describe aspects of a graph:
 - **Adjacency Matrix**
 - **Degree Matrix**
 - **Laplacian Matrix**: difference between degree and adjacency matrix
- Then get a good idea of how to partition graph from eigenvalues and eigenvectors of its Laplacian matrix

Recall: Eigenvalues and Eigenvectors

- The transformation matrix $A = \begin{bmatrix} 2 & 1 \\ 1 & 2 \end{bmatrix}$ preserves the direction of vectors parallel to $\mathbf{v} = (1, -1)^T$ (in purple) and $\mathbf{w} = (1, 1)^T$ (in blue). The vectors in red are not parallel to either eigenvector, so, their directions are changed by the transformation.

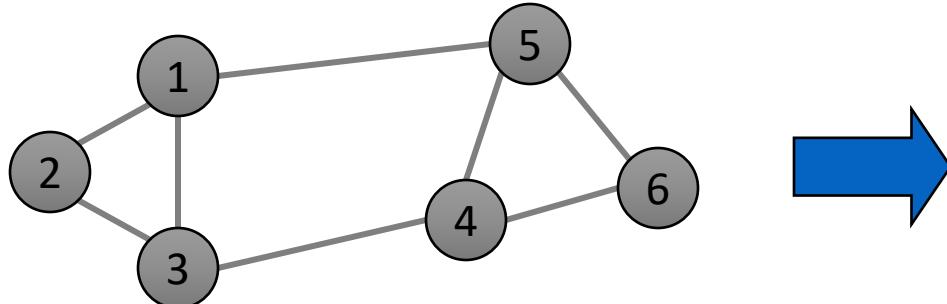
$$A\mathbf{v} = \lambda\mathbf{v}$$

<http://setosa.io/ev/eigenvectors-and-eigenvalues/>
https://en.wikipedia.org/wiki/Eigenvalues_and_eigenvectors



Matrix Representations

- **Adjacency matrix (A):**
 - $n \times n$ matrix
 - $A=[a_{ij}]$, $a_{ij}=1$ if edge between node i and j



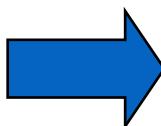
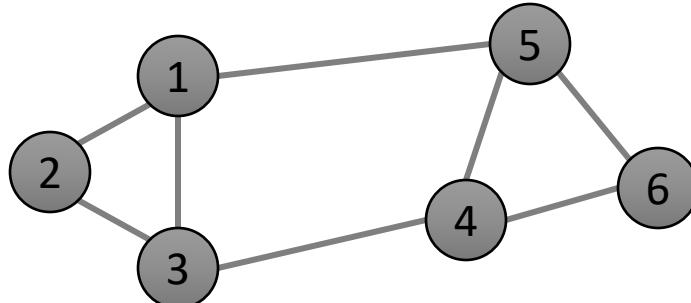
	1	2	3	4	5	6
1	0	1	1	0	1	0
2	1	0	1	0	0	0
3	1	1	0	1	0	0
4	0	0	1	0	1	1
5	1	0	0	1	0	1
6	0	0	0	1	1	0

- **Important properties:**
 - Symmetric matrix
 - Eigenvectors are real and orthogonal
 - $\text{dot_product}(\text{Eigenvectors}_i, \text{Eigenvectors}_j) = 0$

Matrix Representations

- **Degree matrix (D):**

- $n \times n$ diagonal matrix
- $D=[d_{ii}]$, d_{ii} = degree of node i

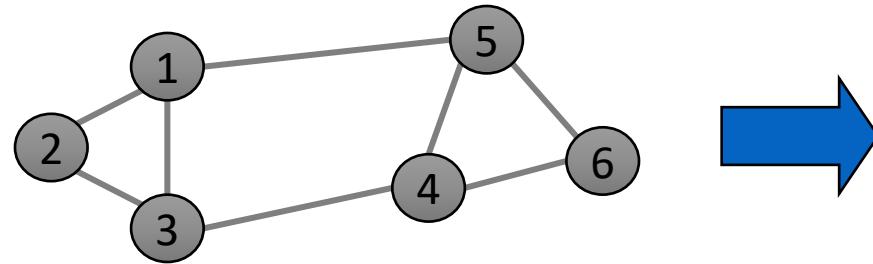


	1	2	3	4	5	6
1	3	0	0	0	0	0
2	0	2	0	0	0	0
3	0	0	3	0	0	0
4	0	0	0	3	0	0
5	0	0	0	0	3	0
6	0	0	0	0	0	2

Matrix Representations

- **Laplacian matrix (L):**

- $n \times n$ symmetric matrix



	1	2	3	4	5	6	
1	3	-1	-1	0	-1	0	
2	-1	2	-1	0	0	0	
3	-1	-1	3	-1	0	0	
4	0	0	-1	3	-1	-1	
5	-1	0	0	-1	3	-1	
6	0	0	0	-1	-1	2	

- **What is trivial eigenpair?**

- $x = (1, \dots, 1)$ then $L \cdot x = 0$ and so $\lambda = \lambda_1 = 0$ (smallest eigenvalue)

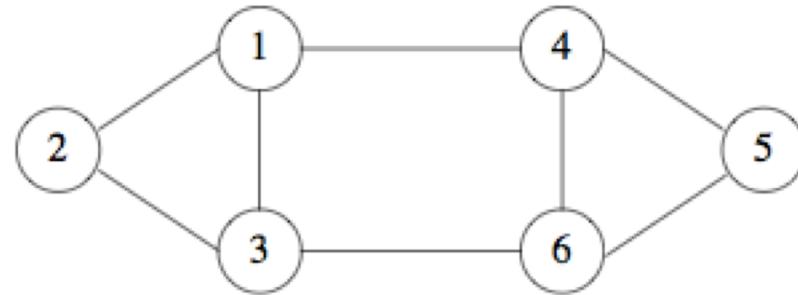
- **Important properties of symmetric matrices:**

- **Eigenvalues** are non-negative real numbers
 - **Eigenvectors** are real and orthogonal

$$L = D - A$$

$$\mathbf{x}^T \mathbf{1} = \sum_{i=1}^n x_i = 0$$

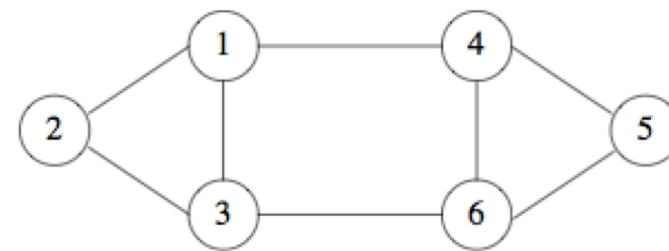
Example 10.19



$$\begin{bmatrix} 3 & -1 & -1 & -1 & 0 & 0 \\ -1 & 2 & -1 & 0 & 0 & 0 \\ -1 & -1 & 3 & 0 & 0 & -1 \\ -1 & 0 & 0 & 3 & -1 & -1 \\ 0 & 0 & 0 & -1 & 2 & -1 \\ 0 & 0 & -1 & -1 & -1 & 3 \end{bmatrix}$$

- Graph and its Laplacian matrix

Example (cont.)



Eigenvalue	0	1	3	3	4	5
Eigenvector	1	1	-5	-1	-1	-1
	1	2	4	-2	1	0
	1	1	1	3	-1	1
	1	-1	-5	-1	1	1
	1	-2	4	-2	-1	0
	1	-1	1	3	1	-1

- Use standard math package to find all eigenvalues and eigenvectors
 - (Have not scaled eigenvectors to length 1, but could)
- Second eigenvector has three positive and three negative components
- Suggest obvious partitioning of {1,2,3} and {4,5,6}

λ_2 as optimization problem

- Recall: to find **second-smallest eigenvalue** for symmetric matrix (such as Laplacian):
 - Second smallest eigenvalue is the minimum of $x^T L x$ where $x = [x_1, x_2, \dots, x_n]$ is a column vector ($n \times 1$, $n = \#$ of nodes) (Rayleigh quotient)
 - Sum of $x_i^2 = 1$
 - x is orthogonal to the eigenvector associated with smallest eigenvalue
- **What is the meaning of $\min x^T L x$ on G ?**
 - $x^T L x = \sum_{i,j=1}^n L_{ij} x_i x_j = \sum_{i,j=1}^n (D_{ij} - A_{ij}) x_i x_j$
 - $= \sum_i d_i x_i^2 - \sum_{(i,j) \in E} 2x_i x_j$
 - $= \sum_{(i,j) \in E} (x_i^2 + x_j^2 - 2x_i x_j) = \sum_{(i,j) \in E} (\underline{x_i - x_j})^2$



Node i has degree d_i . So, value x_i^2 needs to be summed up d_i times.
But each edge (i, j) has two endpoints so we need $x_i^2 + x_j^2$

λ_2 as optimization problem (cont'd)

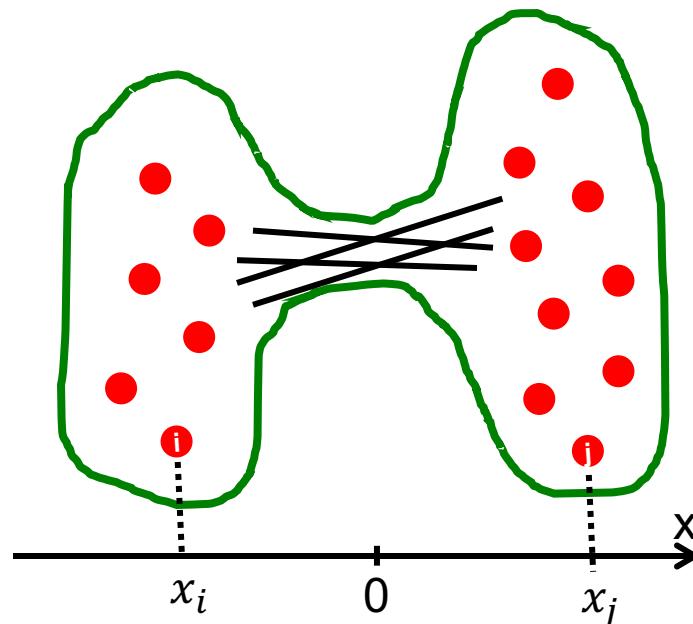
- What else do we know about x ?

- x is unit vector: $\sum_i x_i^2 = 1$
- x is orthogonal to 1st eigenvector $(1, \dots, 1)$ thus:
 $\sum_i x_i \cdot 1 = \sum_i x_i = 0$

- Remember:

$$\lambda_2 = \min_{\substack{\text{All labelings} \\ \text{of nodes } i \text{ so} \\ \text{that } \sum x_i = 0}} \frac{\sum_{(i,j) \in E} (x_i - x_j)^2}{\sum_i x_i^2}$$

We want to assign values x_i to nodes i such
that few edges cross 0.
(we want x_i and x_j to subtract each other)



Ncut as an optimization problem

$$\min_{\mathbf{x}} Ncut(\mathbf{x}) = \min_{\mathbf{y}} \frac{\mathbf{y}^T (\mathbf{D} - \mathbf{W}) \mathbf{y}}{\mathbf{y}^T \mathbf{D} \mathbf{y}},$$

$$\mathbf{y} = (\mathbf{1} + \mathbf{x}) - b(\mathbf{1} - \mathbf{x}),$$

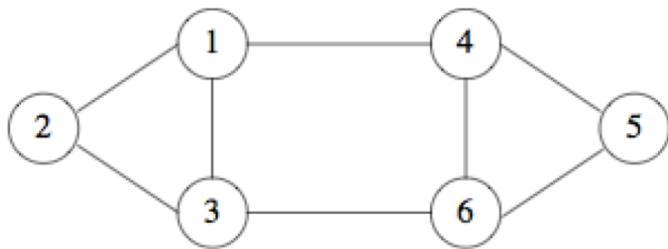
$$b = \frac{k}{1-k}.$$

Given a partition of nodes of a graph, V , into two sets A and B, let \mathbf{x} be an $N = |V|$ dimensional indicator vector, $x_i = 1$ if node i is in A and -1 , otherwise. Let $\mathbf{d}(i) = \sum_j w(i, j)$ be the

Now, recall a simple fact about the *Rayleigh quotient* [11]:

Let \mathbf{A} be a real symmetric matrix. Under the constraint that \mathbf{x} is orthogonal to the $j-1$ smallest eigenvectors $\mathbf{x}_1, \dots, \mathbf{x}_{j-1}$, the quotient $\frac{\mathbf{x}^T \mathbf{A} \mathbf{x}}{\mathbf{x}^T \mathbf{x}}$ is minimized by the next smallest eigenvector \mathbf{x}_j and its minimum value is the corresponding eigenvalue λ_j .

Recall: Example



Eigenvalue	0	1	3	3	4	5
Eigenvector	1	1	-5	-1	-1	-1
	1	2	4	-2	1	0
	1	1	1	3	-1	1
	1	-1	-5	-1	1	1
	1	-2	4	-2	-1	0
	1	-1	1	3	1	-1

- Use standard math package to find all eigenvalues and eigenvectors
 - (Have not scaled eigenvectors to length 1, but could)
- Second eigenvector has three positive and three negative components
- Suggest obvious partitioning of {1,2,3} and {4,5,6}

So far...

- **How to define a “good” partition of a graph?**
 - Minimize a given graph cut criterion
- **How to efficiently identify such a partition?**
 - Approximate using information provided by the eigenvalues and eigenvectors of a graph
- **Spectral Clustering**
 - Naïve approach:
 - Split at 0

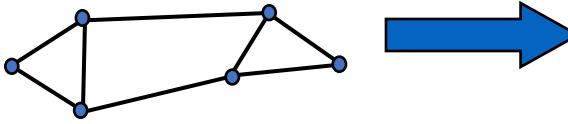
Spectral Clustering Algorithms

- Three basic stages:
 - 1) **Pre-processing**
 - Construct a matrix representation of the graph
 - 2) **Decomposition**
 - Compute eigenvalues and eigenvectors of the matrix
 - Map each point to a lower-dimensional representation based on one or more eigenvectors
 - 3) **Grouping**
 - Assign points to two or more clusters, based on the new representation

Spectral Partitioning Algorithm

- 1) Pre-processing:

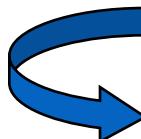
- Build Laplacian matrix L of the graph



	1	2	3	4	5	6
1	3	-1	-1	0	-1	0
2	-1	2	-1	0	0	0
3	-1	-1	3	-1	0	0
4	0	0	-1	3	-1	-1
5	-1	0	0	-1	3	-1
6	0	0	0	-1	-1	2

- 2) Decomposition:

- Find eigenvalues λ and eigenvectors x of the matrix L



0.0
1.0
3.0
3.0
4.0
5.0

0.4	0.3	-0.5	-0.2	-0.4	-0.5
0.4	0.6	0.4	-0.4	0.4	0.0
0.4	0.3	0.1	0.6	-0.4	0.5
0.4	-0.3	0.1	0.6	0.4	-0.5
0.4	-0.3	-0.5	-0.2	0.4	0.5
0.4	-0.6	0.4	-0.4	-0.4	0.0

- Map vertices to corresponding components of λ_2

1	0.3
2	0.6
3	0.3
4	-0.3
5	-0.3
6	-0.6

How do we now find the clusters?

Spectral Partitioning

- 3) **Grouping:**

- Sort components of reduced 1-dimensional vector
- Identify clusters by splitting the sorted vector in two

- **How to choose a splitting point?**

- Naïve approaches:

- Split at **0** or median value

- More expensive approaches:

- Attempt to minimize normalized cut in 1-dimension
(sweep over ordering of nodes induced by the eigenvector)

Split at 0:

Cluster A: Positive points

Cluster B: Negative points

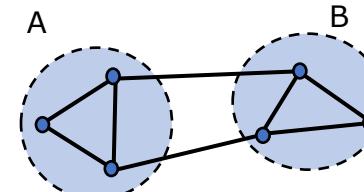


1	0.3
2	0.6
3	0.3
4	-0.3
5	-0.3
6	-0.6

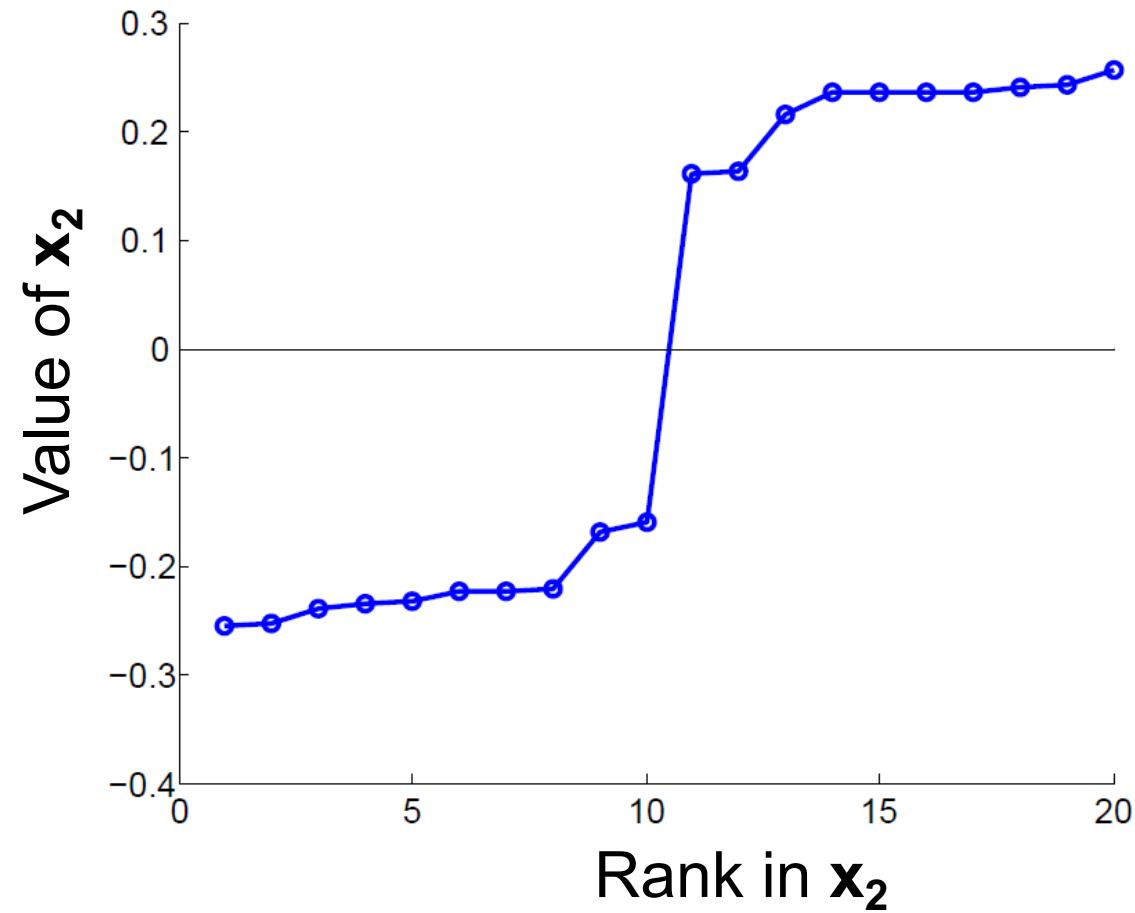
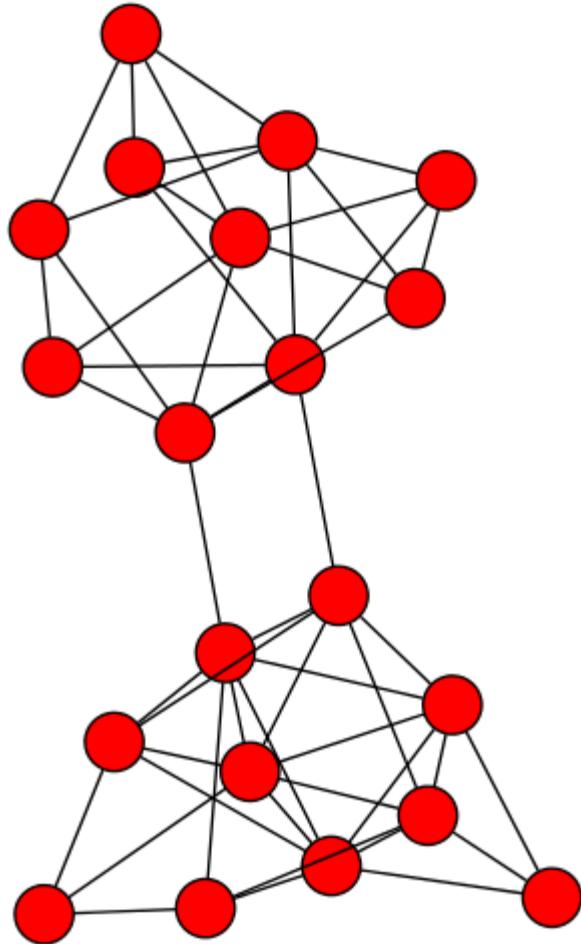


1	0.3
2	0.6
3	0.3

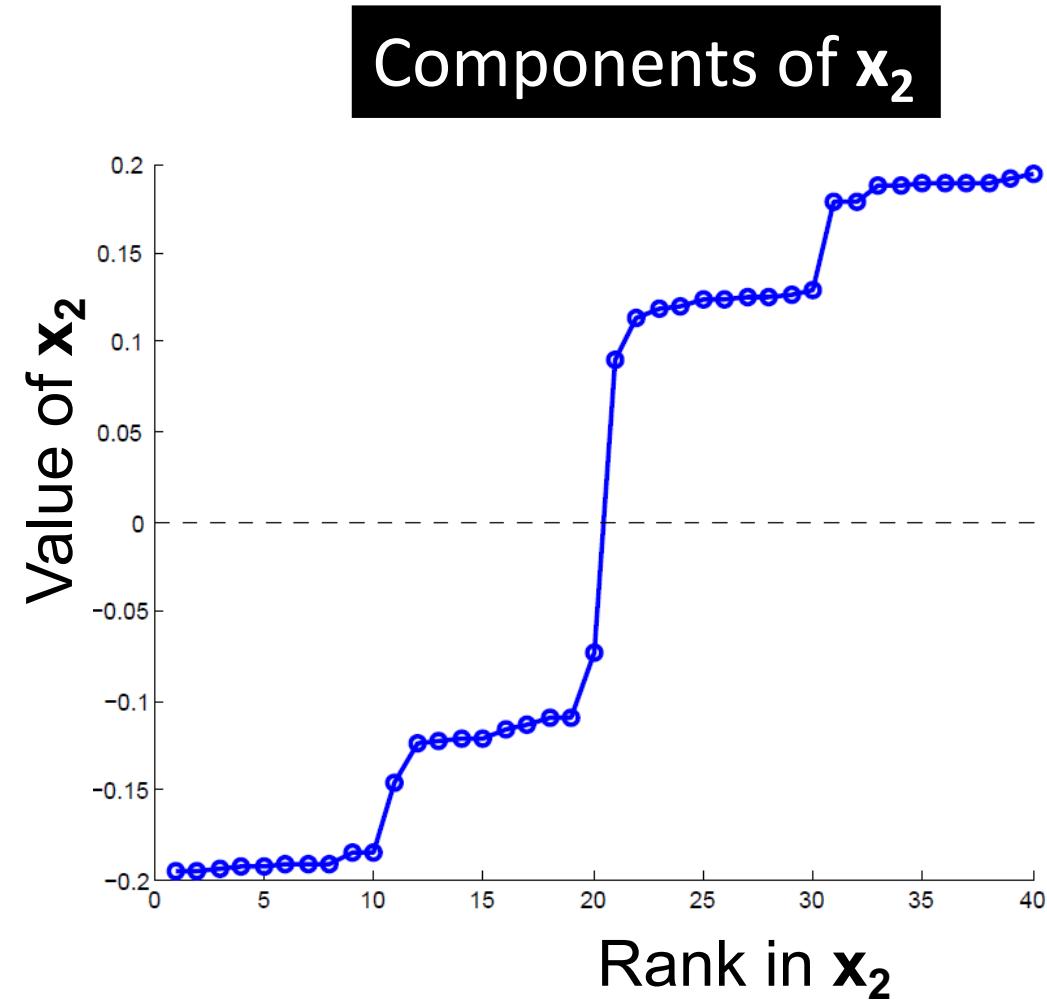
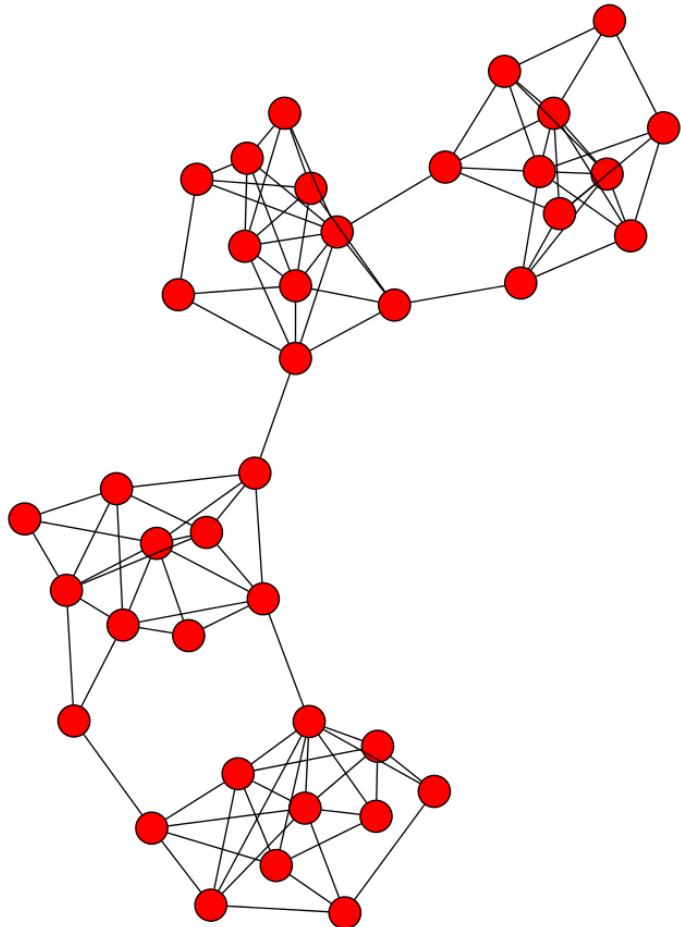
4	-0.3
5	-0.3
6	-0.6



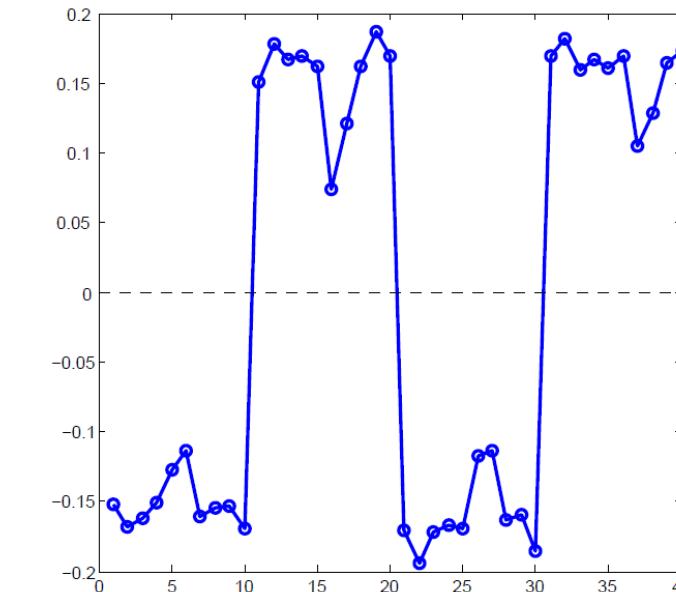
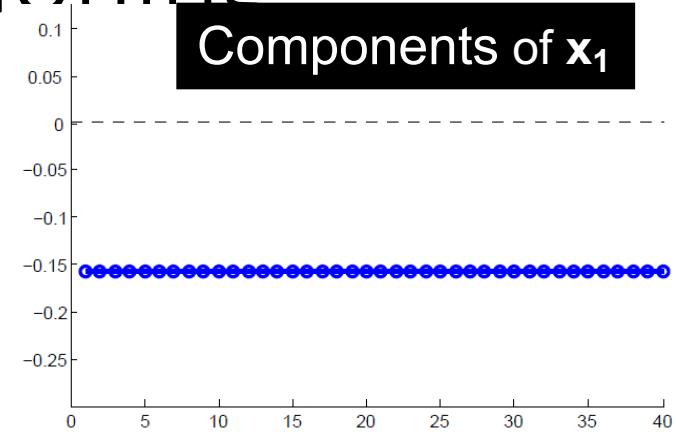
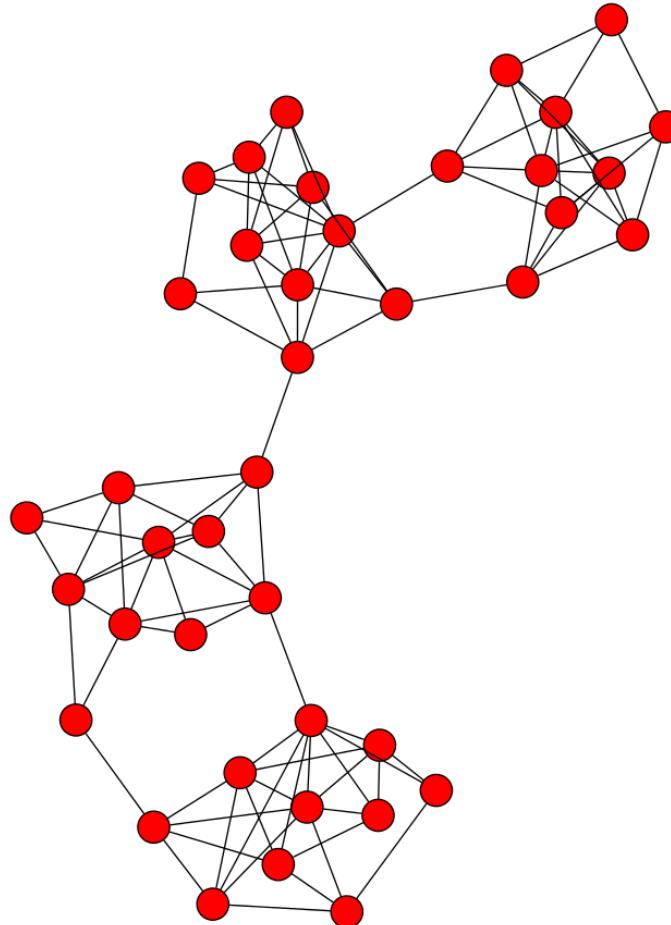
Example: Spectral Partitioning



Example: Spectral Partitioning



Example: Spectral partitioning



Components of \mathbf{x}_3

k-Way Spectral Clustering

- How do we partition a graph into k clusters?
- Two basic approaches:
 - Recursive bi-partitioning [Hagen et al., '92]
 - Recursively apply bi-partitioning algorithm in a hierarchical divisive manner
 - Disadvantages: Inefficient, unstable
 - Cluster multiple eigenvectors [Shi-Malik, '00]
 - Build a reduced space from multiple eigenvectors
 - Commonly used in recent papers
 - Multiple eigenvectors prevent instability due to information loss
 - A preferable approach...